

# Axiomatic Systems Programming: Structural Type Theory and Exact Arithmetic in C++23

The `dedekind` Authors

April 10, 2026

## Abstract

This paper introduces `dedekind`, a C++23 library designed to embed formal mathematical structures directly within the language’s type system as first-class, verified specifications. We address the performance-productivity gap in computational science—often characterized as the “two-language problem”—by leveraging modern C++ features, including modules, concepts, and non-type template parameters (NTTP), to implement a structural type theory. Unlike traditional object-oriented frameworks that rely on nominal inheritance, `dedekind` reifies mathematical entities as positions within a categorical system of relations.

We demonstrate the utility of this approach through a formal construction of the continuum  $\mathbb{R}$  via Dedekind cuts, enabling the exact resolution of transcendental numbers without floating-point overhead. Furthermore, we show how hoisting mereological and algebraic invariants into type signatures allows the LLVM backend to perform aggressive structural pruning, such as collapsing contradictory set-builder predicates into zero-instruction machine code. By addressing language-level constraints through the introduction of the “Highway Notation” for Kleisli composition, this work illustrates that integrating formal mathematical rigor into systems programming can yield zero-overhead abstractions suitable for high-performance symbolic computing. We propose this methodology as a nascent paradigm for *Axiomatic Systems Programming*, where the compiler serves as a formal verification engine for mathematical truth.

## Contents

|          |                                                                                    |          |
|----------|------------------------------------------------------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                                                                | <b>2</b> |
| <b>2</b> | <b>Implementation</b>                                                              | <b>5</b> |
| 2.1      | Instrumented Model Validation: The CI Pipeline as an Independent Auditor . . .     | 5        |
| 2.1.1    | Nominal vs. Structural Subtyping . . . . .                                         | 6        |
| 2.1.2    | The Compiler as a Verification and Optimization Engine . . . . .                   | 7        |
| 2.1.3    | Ontological Decidability and the Build Graph . . . . .                             | 7        |
| 2.2      | Level 0: The Algebraic Atlas ( <code>dedekind.category</code> ) . . . . .          | 8        |
| 2.2.1    | The <code>:species</code> partition: The Algebraic Atlas . . . . .                 | 8        |
| 2.2.2    | The <code>:morphism</code> partition: The Skeletal Arrow . . . . .                 | 9        |
| 2.3      | The <code>:morphism</code> partition: The Skeletal Arrow . . . . .                 | 11       |
| 2.3.1    | The <code>:mereology</code> partition: The Language of Parts . . . . .             | 11       |
| 2.3.2    | The <code>:logic</code> partition: The Subobject Classifier ( $\Omega$ ) . . . . . | 12       |
| 2.3.3    | Posetal Categories ( <code>:posetal</code> ) . . . . .                             | 13       |
| 2.3.4    | Small Categories ( <code>:small</code> ) . . . . .                                 | 14       |
| 2.3.5    | Universal Products and Coproducts ( <code>:cartesian</code> ) . . . . .            | 15       |
| 2.3.6    | Boundary Objects ( <code>:limit</code> ) . . . . .                                 | 16       |
| 2.3.7    | Universal Constructions and Pullbacks ( <code>:pullback</code> ) . . . . .         | 16       |
| 2.4      | Level 1: Space and Action ( <code>dedekind.category</code> ) . . . . .             | 17       |

|          |                                                                                |           |
|----------|--------------------------------------------------------------------------------|-----------|
| 2.4.1    | The <code>:total</code> partition: The Property of Totality . . . . .          | 17        |
| 2.4.2    | The <code>:algebra_total</code> partition: The Laws of Harmony . . . . .       | 17        |
| 2.4.3    | Monoid Actions and Modules ( <code>:actions</code> ) . . . . .                 | 18        |
| 2.4.4    | The <code>:functor</code> partition: Mappings between Categories . . . . .     | 19        |
| 2.4.5    | The <code>:natural</code> partition: Natural Transformations . . . . .         | 19        |
| 2.4.6    | The <code>:monad</code> partition: The $\Omega$ -Monad . . . . .               | 20        |
| 2.4.7    | The <code>:kleisli</code> partition: The Universal Highway . . . . .           | 20        |
| 2.4.8    | The <code>:partial</code> partition: The Pragmatic Rescue . . . . .            | 21        |
| 2.4.9    | The <code>:algebra_partial</code> partition: The Generalised Reality . . . . . | 21        |
| 2.4.10   | The <code>:numeric</code> partition: The Material Landing . . . . .            | 22        |
| 2.4.11   | The <code>:etcs</code> partition: The Set-Theoretic Interface . . . . .        | 22        |
| 2.5      | <code>sets</code> Module . . . . .                                             | 23        |
| 2.5.1    | The <code>:set</code> partition . . . . .                                      | 23        |
| 2.5.2    | Intensional Sets and Computational Undecidability . . . . .                    | 24        |
| 2.5.3    | Categorification of Sequences and the Continuum . . . . .                      | 24        |
| <b>3</b> | <b>Results</b> . . . . .                                                       | <b>27</b> |
| 3.1      | The Registry of Formal Correspondences . . . . .                               | 27        |
| 3.2      | <code>order</code> Module . . . . .                                            | 29        |
| 3.3      | <code>topology</code> Module . . . . .                                         | 29        |
| 3.4      | <code>sequences</code> Module . . . . .                                        | 29        |
| 3.5      | <code>algebra</code> Module . . . . .                                          | 29        |
| 3.6      | <code>morphologies</code> Module . . . . .                                     | 29        |
| 3.7      | <code>geometry</code> Module . . . . .                                         | 29        |
| 3.8      | <code>numbers</code> Module . . . . .                                          | 29        |
| 3.9      | <code>analysis</code> Module . . . . .                                         | 29        |
| <b>4</b> | <b>Discussion</b> . . . . .                                                    | <b>29</b> |
| <b>5</b> | <b>Conclusion</b> . . . . .                                                    | <b>30</b> |

## 1 Introduction

The development of modern computational science, particularly in data science and Large Language Model (LLM) engineering, is constrained by a persistent performance-productivity gap known as the *two-language problem*. In this paradigm, researchers utilize high-level symbolic environments like `Python` for rapid prototyping and expressive reasoning, while relying on low-level languages like `C++` to handle execution backends. While tools such as `PyTorch` [1] and `NumPy` [2] offer elegant interfaces, they typically treat mathematical logic as a runtime implementation detail. This decoupling forces a manual translation layer between the high-level mathematical intent and the underlying hardware execution, leading to significant overhead in both developer time and system latency.

This duality is especially relevant in the context of LLM development and large-scale symbolic computing. In these fields, the ability to maintain the structural integrity of a mathematical model—such as the exactness of a tensor operation or the logical consistency of a set-builder predicate—is often lost during the transition to the `C++` backend. Traditional integration methods, such as `ctypes` or even modern binding tools like `nanobind` [3], bridge the binary gap but do not bridge the *semantic* gap. They allow `Python` to call `C++` functions, but they do not allow the `C++` compiler to reason about the high-level mathematical invariants defined in the symbolic layer.

Recent advances in systems programming have begun to challenge the historical divide between low-level performance and high-level abstraction. The `Rust` programming language, for

instance, has gained significant traction by providing a "safety-first" approach through its affine type system and trait-based polymorphism. However, while Rust excels at memory safety and nominal constraints, its type system often operates as a "black box" that can obscure the structural transparency required to hoist complex mathematical invariants into the heart of the optimization pipeline.

In contrast, the evolution of C++ toward the C++23 standard [4] has introduced a suite of features—notably `concepts`, `constexpr`, and non-type template parameters (NTTP)—that enable a form of *structural typing*. Unlike traditional object-oriented hierarchies, these tools allow developers to define "mathematical species" based on their internal requirements and relational properties rather than their inherited names. As explored by Čukić [5], modern C++ is now sufficiently advanced to simulate, and in some cases surpass, the expressive power of pure functional programming languages. By treating types not merely as memory layouts but as compile-time predicates, C++ can now reify abstract "notions of computation" directly into the machine code.

This shift allows for the embedding of category-theoretic patterns, such as monads and functors, which serve as the formal foundation for structured side-effects and composition [6]. In the `dedekind` library, we leverage this functional substrate to treat mathematical entities as positions in a system of relations. Specifically, we utilize C++ `concepts` to enforce structural integrity at the call site, ensuring that the compiler acts as a verification engine. This approach transforms the C++ backend from a simple execution target into a sophisticated environment capable of reasoning about symbolic logic—effectively turning the "slow oracle" of a high-level symbolic system into a zero-overhead, statically-verified binary.

This shift toward compile-time verification mirrors a long-standing debate in the philosophy of mathematics between substance-based and structuralist ontologies. While traditional object-oriented programming relies on *nominal typing*—where an object's identity is tied to its named inheritance from a specific class—modern C++ features like `concepts` facilitate a *structural* approach. This aligns with the mathematical structuralism championed by the Bourbaki group and further formalized through category theory [7]. In this view, mathematical entities are not defined by what they "are" in isolation, but by the system of relations they inhabit.

The `dedekind` library specifically adopts the framework of *combinatorial species* introduced by Joyal [8, 9]. By treating mathematical structures as "species"—mappings that describe how a structure is placed upon a set of labels—we can reify abstract species (such as rings, fields, or even the continuum itself) as C++ template specifications. This "white-box" structural transparency allows the compiler to reason about the properties of a type (e.g., its commutativity or its completeness) purely by inspecting its requirements, much like McLarty's observation that "numbers can be just what they have to" within a categorical framework [10].

```
template <typename Universe, auto Predicate>
struct Set {
    using type = Universe;
    static constexpr auto p = Predicate;
};

// Mereological composition (Intersection)
template <typename S1, typename S2>
using Intersection = Set<typename S1::type, [] (auto x) {
    return S1::p(x) && S2::p(x);
}>;
```

Listing 1: A simplified illustration of a set intersection implementation in C++ in the structuralist, mereological style.

In practice, this allows us to move beyond the "black-box" constraints of nominal traits. For instance, while a nominal system might require a type to explicitly `impl Real`, a structural

system in C++ can identify a type as a representation of  $\mathbb{R}$  if it satisfies the structural requirements of a Dedekind-complete ordered field. By hoisting these mereological invariants into the type signature, we enable the LLVM backend to perform aggressive structural pruning. As noted in the *n-Category Café* [11], this categorical perspective treats the "logic" of the system and the "computation" of the system as dual aspects of the same structure—a principle we aim to exploit. For example, to ensure that contradictory set-builder predicates collapse into zero-instruction machine code (see Listing 1).

```
using
  = Set<
  , [](auto x) { return x > 0.0; }>;
using
  = Set<
  , [](auto x) { return false; }>;
// Intersection with
  is symbolically reduced
using Result = Intersection<
  ,
>;
// Proven at compile-time: Result is the empty set
static_assert(Result::p(5.0) == false);
```

Listing 2: A simplified illustration of the compiler simplifies the intersection with the empty set to a zero-cost identity.

The utility of this structuralist approach is best illustrated through two distinct computational advantages. First, the library enables aggressive structural pruning at compile-time. By reifying set operations as type-level compositions, the compiler can symbolically resolve the relationship between predicates. For instance, the intersection of two contradictory sets—such as the set of positive reals and the set of negative reals—can be identified as an empty set during translation to LLVM Intermediate Representation (IR). This allows the backend to prune entire execution branches, effectively collapsing complex symbolic logic into zero-instruction machine code.

```
// Representing e via the series predicate:
// q < sum_{n=0}^N 1/n! for some N
using e = Set<Q, [](Q q) {
  // Symbolic check: ∃ N such that q < partial_sum(N)
  // In dedekind, this is resolved by the compiler's
  // ability to evaluate monotonic convergence.
  return ∃_approximation(q, [](int n) {
    return factorial_reciprocal_sum(n);
  });
}>;

// The compiler resolves this comparison at translation time
static_assert(e::p({271, 100}) == true); // 2.71 < e
```

Listing 3: Illustration of the symbolic construction of  $e$  as an intensional set (the lower cut  $L \subset \mathbb{Q}$ ).

Furthermore, **dedekind** provides the ability to model intensional, transfinite sets through a purely symbolic approach. Rather than representing a collection as an extensional buffer of values, we define it by its membership condition. This is most powerfully demonstrated in our construction of the continuum  $\mathbb{R}$  via Dedekind cuts. By representing a real number as a compile-time coordinate—a partition of the rational numbers  $\mathbb{Q}$ —we enable the exact resolution of transcendental values without the rounding errors or overhead of standard floating-point

arithmetic [12]. As shown in Listing 4, comparing  $e$  and  $\pi$  becomes a compile-time verification of their respective partitions.

```
// Defined as the lower cut of their respective series/properties
using e    = Set<Q, [](auto q) { return  $\exists$ _lower_bound(q, e_series); }>;
using  $\pi$     = Set<Q, [](auto q) { return  $\exists$ _lower_bound(q,  $\pi$ _series); }>;

/**
 * Structural comparison: Is e < pi?
 * This is true if there exists a rational 'r' such that:
 * r is in the lower cut of pi, but NOT in the lower cut of e.
 */
static_assert(e <  $\pi$ );
```

Listing 4: Illustration of the comparison of the transcendental constants  $e$  and  $\pi$  via their Dedekind partitions.

Together, these features bridge the gap between the declarative ease of symbolic algebra and the mechanical sympathy required for modern high-performance computing [13]. By allowing the C++ compiler itself to understand the symbolic logic of the Continuum, **dedekind** ensures that high-level mathematical invariants are preserved from specification to machine code, effectively treating the compiler as a formal proof assistant for systems-level performance.

## 2 Implementation

The methodology of **dedekind** is predicated on the translation of formal mathematical structures into a stratified registry of C++23 modules. This stratification allows for a "bottom-up" synthesis of complex mathematical objects, where each level serves as a verified foundation for the next. By design, the flow of compilation mirrors a pedagogical progression: foundational, general-purpose modules are verified first, providing an immutable basis for the specialized and practical abstractions that follow.

Table 1: The Dedekind Atomic Floor: Primary Formal Correspondences

| Mathematical Concept               | C++23 Implementation              | Structural concept            |
|------------------------------------|-----------------------------------|-------------------------------|
| <i>Morphic Foundations</i>         |                                   |                               |
| Morphism ( $f : A \rightarrow B$ ) | <code>std::invoke_result_t</code> | <code>IsArrow</code>          |
| Terminal Object (1)                | <code>std::monostate</code>       | <code>IsTerminalObject</code> |
| Initial Object (0)                 | <code>std::nullptr_t</code>       | <code>IsInitialObject</code>  |
| <i>Structural Relations</i>        |                                   |                               |
| Parthood ( $\sqsubseteq$ )         | <code>operator&lt;=</code>        | <code>IsPartOf</code>         |
| Subobject Classifier ( $\Omega$ )  | <code>enum class Omega</code>     | <code>LogicalSpecies</code>   |
| <i>Universal Constructions</i>     |                                   |                               |
| Product ( $A \times B$ )           | <code>std::pair</code>            | <code>IsProduct</code>        |
| Coproduct ( $A + B$ )              | <code>std::variant</code>         | <code>IsCoproduct</code>      |

### 2.1 Instrumented Model Validation: The CI Pipeline as an Independent Auditor

While formal verification in languages like Agda or Coq typically relies on static type-level proofs that are erased during compilation, the **dedekind** framework adopts a methodology of *Instrumented Model Validation*. By leveraging the LLVM profiling infrastructure, we provide an empirical *Runtime Witness* for our categorical primitives.

This approach serves as a deterministic evolution of the property-based testing tradition pioneered by QUICKCHECK [14]. Where QUICKCHECK utilizes stochastic search to falsify algebraic properties in Haskell, `dedekind` utilizes the CI/CD pipeline, which exercises `static_assert` (evaluated at compile-time) combined with `Catch2` tests (executed after unit compilation) to ensure that every structural identity is physically materialized in the binary. By enforcing high levels of patch coverage on new ontological partitions, we provide a "Physical Proof" in the form of *Binary Module Interface* (BMI) files that the abstract category theory—often treated as a ghost in the machine—is actually traversed by the hardware. This aligns with the "Real World" functional programming philosophy advocated by O’Sullivan et al. [15], where formal rigor is balanced with the practical necessity of observable, high-performance execution.

This methodology is physically enforced by a strictly layered directed acyclic graph (DAG) of the system’s build order. By carefully selecting C++ module partitions and build-step boundaries, we translate the dependencies of formal mathematics into a sequence of verified compilation units. This architecture mimics the pedagogical progression of a mathematical textbook, where the ontology matures from foundational axioms to complex theorems through a series of verified proofs. By isolating these prerequisites into a ‘category’ module, we leverage what Wadler termed "Theorems for Free!" [16]: once a structural identity is established and cached as a Binary Module Interface (BMI), its properties are inherited by all downstream partitions as an immutable "cached truth." This ensures that the compiler’s proof-search is grounded in a physical hierarchy, preventing circular reasoning and ensuring that the structural integrity of the model is a prerequisite for successful linking. Crucially, the CI/CD methodology provides an independent, physical verification of this progression; by requiring every node in the build graph to pass its respective model validation before the pipeline can proceed, we ensure that the "textbook" derivation is not merely a logical possibility, but a materialized reality across the entire module graph.

### 2.1.1 Nominal vs. Structural Subtyping

Traditional object-oriented systems typically rely on *nominal subtyping*, where a type  $T$  is recognized as belonging to a category  $C$  only through explicit inheritance declarations (e.g., `class T : public Monoid`). In this "Capulet/Montague" paradigm, a type is rejected if it lacks the correct "family name," regardless of its actual properties.

In contrast, `dedekind` adopts a structural type theory. This alignment mirrors the mathematical structuralism advocated by Awodey [7], where mathematical objects are defined not by an internal "essence" but by their roles and relations within a category. As famously noted by Shakespeare, "a rose by any other name would smell as sweet" [17]. In `dedekind`, we codify this by prioritizing the "scent" (structural invariants) over the "lineage" (nominal inheritance), as illustrated in Listing 5.

```
// Nominal: The "Feudal" approach (Inheritance-based)
struct Flower { virtual bool smells_sweet() const = 0; };
struct Rose : public Flower {
    bool smells_sweet() const override { return true; }
};

// Structural: The "Juliet" approach (Concept-based)
template <typename T>
concept SmellsSweet = requires(T t) {
    { t.scent() } -> std::same_as<Sweetness>;
};

// In dedekind, even a 'Briar' is a 'Rose' if it satisfies the scent invariant.
static_assert(SmellsSweet<Briar>, "A rose by any other name...");
```

Listing 5: The "Juliet" Posture: Nominal vs. Structural Typing in C++23.

Leveraging C++20/23 Concepts, types are treated as mathematical objects defined entirely by their interface and satisfied invariants. A type  $T$  is identified as a `MonoidalCategory` if it possesses the required algebraic morphisms  $(\otimes, e)$ ; its internal name and declaration lineage are irrelevant to the compiler’s logical resolution.

### 2.1.2 The Compiler as a Verification and Optimization Engine

While functional languages like Haskell are often viewed as the natural home for Category Theory, Milewski [18] demonstrated that C++ is equally capable of modeling complex categorical morphisms. `dedekind` builds upon this bridge by utilizing C++23 features to move from mere simulation to compile-time verification.

By applying `static_assert` against structural concepts, we transform the compilation process into a formal verification phase. This provides the LLVM toolchain with the semantic context required for “structural pruning”: by hoisting mereological relations into the type signature, the backend can identify logical contradictions—such as an empty set defined by contradictory predicates—and prune them into zero-overhead machine instructions. This “correctness-by-construction” methodology ensures that the structural integrity of the mathematical model is verified and optimized before the final binary is generated.

### 2.1.3 Ontological Decidability and the Build Graph

To mitigate the “template tax” and ensure the practical decidability of the ontology, we utilize a directed acyclic build graph (see Figure 2.1.3). This directed acyclic graph mirrors the formal structure of a mathematical proof: *Axiom*  $\rightarrow$  *Theorem*  $\rightarrow$  *Proof*. While the underlying mathematical nature of certain propositions remains inherently undecidable, `dedekind` provides structural guardrails by isolating foundational axioms into up-stream module partitions. The resulting BMI acts as the “cached truth” established in the previous section.

This physical stratification prohibits circular reasoning at the compiler level as a partition can only reference a `concept` which was previously defined. Subsequent partitions then import these BMIs as established, immutable lemmata, allowing the compiler to focus on the synthesis of higher-order laws with a greatly reduced risk of recursive definition or redundant re-instantiation. While the module graph serves as an external meta-logical constraint ensuring build-time termination, we later introduce a pluggable *Subobject Classifier* ( $\Omega$ ) in the ?? section which internalizes the notion of undecidability within the C++ type system itself.




Figure 1: Module build dependencies in the `dedekind` library.

For the remainder of the implementation section, we simply follow build order of the system. We will annotate and document the specific modules and partitions as they are defined and validated, beginning with the foundational internal partitions of the `dedekind.category` module.

In this approach, the structure of the paper is intended to be isomorphic to the CI/CD pipeline: the text serves as a formal documentation of the build, while the build, in turn, provides a physical validation of the claims made within the paper. By the time the reader reaches the final sections, the library’s ‘textbook’ derivation will have been fully materialized as a verified C++ library object traversed by the underlying hardware.

## 2.2 Level 0: The Algebraic Atlas (`dedekind.category`)

While the global build graph establishes the dependency between C++ modules, the `dedekind.category` module constitutes the system's "Atomic Floor." To maintain absolute linear order, this module is itself stratified into internal partitions. These partitions represent the maturation of mathematical species from raw mappings (`:morphism`) to structured universes (`:logic`).




Figure 2: Refined Build Graph: The "Numeric" layer now serves as the base for "Partial" logic, allowing "ETCS" to remain decoupled from hardware specificities.

### 2.2.1 The `:species` partition: The Algebraic Atlas

The `:species` partition serves as the foundational *Algebraic Atlas* of the library, providing a stratified registry of structural invariants for C++ machine types and standard library operators. Following the combinatorial species theory of Joyal [8, 9], we treat fundamental types not merely as passive sets of values, but as active rules for generating mathematical structure. By reifying categorical axioms—such as associativity and commutativity—as compile-time constants, the library transforms the C++ type system into a formal verification engine.

As illustrated in Table ??, the **Feature Cube** acts as a formal ledger mapping algebraic properties to the language's fundamental atoms. This explicit mapping allows the `dedekind` dispatcher to distinguish between idealized algebraic structures (e.g., `unsigned int` modular addition) and the messy inhabitants of the hardware, such as floating-point operations which exhibit numerical non-associativity under the IEEE 754 standard [12]. By anchoring these operators as first-class categorical entities, we enable the compiler to act as a physical guardrail, ensuring that structural integrity is verified before the final binary is generated.



Crucially, this allows for the *Honest Rejection* of hardware species that violate the laws of the intended mathematical space: while `unsigned int` addition satisfies the *Monoid* axioms, it fails as an *Ordered Monoid* at the Lipschitz boundary due to modular wrap-around. Conversely, `signed int` is treated as a *Hazardous Ordered Monoid* where the compiler exploits UB to enforce a false sense of algebraic harmony through assumed monotonicity.

```
/**
 * @brief The Feature Cube: Categorical axioms as compile-time traits.
 * @tparam T The Species (e.g., int, double).
 * @tparam Op The Binary Operation or Relation (e.g., std::plus<>).
 */
template <typename T, typename Op>
struct SpeciesTraits {
    static constexpr bool is_associative = false;
    static constexpr bool is_commutative = false;
    static constexpr bool is_idempotent = false;
};

// Specialization for 'Ideal' Integral Species
template <>
struct SpeciesTraits<int, std::plus<int>> {
    static constexpr bool is_associative = true;
    static constexpr bool is_commutative = true;
};

// Verification: The compiler now "knows" the algebra of the machine.
static_assert(SpeciesTraits<int, std::plus<int>>::is_associative,
    "Integrals form a valid Semigroup.");

static_assert(!SpeciesTraits<double, std::plus<double>>::is_associative,
    "Hardware Floats are non-associative; Rescue required.");
```

Listing 6: The Species Registry: Reifying Algebraic Invariants.

The core objective is to anchor the C++ type system as a substrate for modeling truths obtained from category theory and abstract algebra. In this structuralist setup, providing the axiomatic truths (e.g., that  $a \vee b$  is commutative) or documenting their absence (e.g., *NaN* singularities) becomes the explicit responsibility of the registry. This static grounding provides the necessary "alphabet" for the `:morphism` partition to begin defining the "language" of functional mappings.

## 2.2.2 The `:morphism` partition: The Skeletal Arrow

The `:morphism` partition establishes the primary inhabitant of the `dedekind` topos: the **Arrow** (or **Morphism**). While traditional C++ development treats functions as opaque, executable blocks, this partition reifies them as formal mappings  $f : A \rightarrow B$ . **\*\*In this ontology, the Morphism serves as the "Atom of Composability"\*\***—the fundamental unit from which all higher-level categorical structures are synthesized. By elevating the function object from a procedural instruction to a structural entity, we enable the library to reason about the composition and validity of transformations at compile-time, a strategy central to modern functional C++ architectures [5, 18].

A core mechanism of this partition is **Signature Discovery**. Utilizing C++23 `std::invoke_result_t` and variadic template metaprogramming, the library automatically extracts the domain and codomain of raw lambdas, function pointers, and functors. This reification process registers the callable into the categorical atlas as a first-class inhabitant, providing the "connective tissue" required for the algebraic structures in Level 1 and the functorial lifting in Level 2.

Table 2: The Feature Cube: Hierarchical Axiomatic Mapping. Groups correspond to the **:species** partition. **Laws**: *Identity* ( $x \circ e = x$ ); *Absorber* ( $x \circ z = z$ ); *Associative* ( $(x \circ y) \circ z = x \circ (y \circ z)$ ); *Commutative* ( $x \circ y = y \circ x$ ); *Idempotent* ( $a \circ a = a$ ); *Inverse* ( $a \circ a^{-1} = e$ ); *Monotonic* ( $x < y \implies (x \circ z) < (y \circ z)$ ); *Logic* ( $\chi : A \rightarrow \Omega$ ).

| Species: bool | Logical |   |    | Bitwise |   |   | Logic                                  |  |  |
|---------------|---------|---|----|---------|---|---|----------------------------------------|--|--|
| Operator      |         |   | && |         | & | ^ | $\Omega$                               |  |  |
| Identity      | f       | t |    | f       | t | f | Classical<br>( $\Omega_{\mathbb{B}}$ ) |  |  |
| Absorber      | t       | f |    | t       | f | — |                                        |  |  |
| Associative   | ✓       | ✓ |    | ✓       | ✓ | ✓ |                                        |  |  |
| Commutative   | ✓       | ✓ |    | ✓       | ✓ | ✓ |                                        |  |  |
| Idempotent    | ✓       | ✓ |    | ✓       | ✓ | — |                                        |  |  |
| Inverse       | —       | — |    | —       | — | ✓ |                                        |  |  |
| Monotonic     | ✓       | ✓ |    | ✓       | ✓ | — |                                        |  |  |

| Species: uint | Arithmetic            |                       |                       | Order |     |    | Bitwise |     |   | Logic                                  |
|---------------|-----------------------|-----------------------|-----------------------|-------|-----|----|---------|-----|---|----------------------------------------|
| Operator      | +                     | *                     | %                     | min   | max | <= |         | &   | ^ | $\Omega$                               |
| Identity      | 0                     | 1                     | —                     | ~0u   | 0   | —  | 0       | ~0u | 0 | Classical<br>( $\Omega_{\mathbb{B}}$ ) |
| Absorber      | —                     | 0                     | <b>H</b> <sup>D</sup> | 0     | ~0u | —  | ~0u     | 0   | — |                                        |
| Associative   | ✓ <sup>M</sup>        | ✓ <sup>M</sup>        | —                     | ✓     | ✓   | ✓  | ✓       | ✓   | ✓ |                                        |
| Commutative   | ✓                     | ✓                     | —                     | ✓     | ✓   | —  | ✓       | ✓   | ✓ |                                        |
| Idempotent    | —                     | —                     | —                     | ✓     | ✓   | ✓  | ✓       | ✓   | — |                                        |
| Inverse       | ✓ <sup>M</sup>        | —                     | —                     | —     | —   | —  | —       | —   | ✓ |                                        |
| Monotonic     | <b>H</b> <sup>W</sup> | <b>H</b> <sup>W</sup> | —                     | ✓     | ✓   | ✓  | —       | —   | — |                                        |

| Species: int | Arithmetic            |                |                       | Order |       |    | Bitwise |    |   | Logic                                  |
|--------------|-----------------------|----------------|-----------------------|-------|-------|----|---------|----|---|----------------------------------------|
| Operator     | +                     | *              | %                     | min   | max   | <= |         | &  | ^ | $\Omega$                               |
| Identity     | 0                     | 1              | —                     | I_MAX | I_MIN | —  | 0       | -1 | 0 | Classical<br>( $\Omega_{\mathbb{B}}$ ) |
| Absorber     | —                     | 0              | <b>H</b> <sup>D</sup> | I_MIN | I_MAX | —  | -1      | 0  | — |                                        |
| Associative  | <b>H</b> <sup>U</sup> | ✓              | —                     | ✓     | ✓     | ✓  | ✓       | ✓  | ✓ |                                        |
| Commutative  | ✓                     | ✓              | —                     | ✓     | ✓     | —  | ✓       | ✓  | ✓ |                                        |
| Idempotent   | —                     | —              | —                     | ✓     | ✓     | ✓  | ✓       | ✓  | — |                                        |
| Inverse      | <b>H</b> <sup>A</sup> | —              | —                     | —     | —     | —  | —       | —  | ✓ |                                        |
| Monotonic    | ✓ <sup>P</sup>        | ✓ <sup>P</sup> | —                     | ✓     | ✓     | ✓  | —       | —  | — |                                        |

| Species: double | Arithmetic |     | Order |      |    | Logic                        |  |  |
|-----------------|------------|-----|-------|------|----|------------------------------|--|--|
| Operator        | +          | *   | min   | max  | <= | $\Omega$                     |  |  |
| Identity        | 0.0        | 1.0 | Inf   | -Inf | —  | Ternary<br>( $\Omega_{K3}$ ) |  |  |
| Absorber        | Inf        | NaN | -Inf  | Inf  | —  |                              |  |  |
| Associative     | *          | *   | ✓     | ✓    | ✓  |                              |  |  |
| Commutative     | ✓          | ✓   | ✓     | ✓    | —  |                              |  |  |
| Idempotent      | —          | —   | ✓     | ✓    | ✓  |                              |  |  |
| Inverse         | ✓          | ✓   | —     | —    | —  |                              |  |  |
| Monotonic       | ✓          | ✓   | ✓     | ✓    | ✓  |                              |  |  |

<sup>M</sup> Modular success. <sup>A</sup> Asymmetry hazard (INT\_MIN). <sup>W</sup> Wrap-around hazard.

<sup>U</sup> Undefined behavior (overflow). <sup>D</sup> Division hazard. <sup>P</sup> Partial success.

\*Fails bit-perfect associativity due to rounding/NaN.

As shown in Listing 7, the `IsArrow` concept acts as a formal gatekeeper. It ensures that every mapping is well-formed—possessing a verifiable signature and satisfying the axioms of a morphism—before it can be used as a taxonomic brick in the broader structural hierarchy [19].

```

/** @brief Automatic inference of Morphism metadata f: Domain -> Codomain */
template <typename F, typename... Args>
struct SpeciesTraits<F, Args...> {
    using Domain      = std::tuple_element_t<0, std::tuple<Args...>>;
    using Codomain    = std::invoke_result_t<F, Args...>;
    static constexpr bool is_morphism = true;
};

/** @brief Concept enforcing the integrity of the Skeletal Arrow */
template <typename F, typename... Args>
concept IsArrow = SpeciesTraits<F, Args...>::is_morphism &&
                  requires { typename SpeciesTraits<F, Args...>::Codomain; };

// Reification of a raw lambda into an 'Atomic Arrow'
auto logic_gate = [](int x) -> bool { return x > 0; };

static_assert(IsArrow<decltype(logic_gate), int>,
              "Morphism␣Validated:␣The␣Atom␣of␣Composability␣is␣discovered.");

```

Listing 7: The Atom of Composability: Morphic Reification.

## 2.3 The `:morphism` partition: The Skeletal Arrow

Following the mapping established in Table 1, the `:morphism` partition reifies C++ callables as formal arrows  $f : A \rightarrow B$ . The core mechanism of this partition is **Signature Discovery**, which utilizes `std::invoke_result_t` and variadic template metaprogramming to extract domain and codomain metadata directly from the type system [5, 18].

```

template <typename F, typename... Args>
concept IsArrow = requires {
    typename std::invoke_result_t<F, Args...>; // Codomain exists
    requires sizeof...(Args) > 0; // Domain is non-empty
};

```

Listing 8: The Morphic Reification Hook (Compressed).

This “Atomic Arrow” serves as the fundamental unit from which all higher-level categorical structures are synthesized. By elevating the function object from an opaque procedural block to a structural entity, the library enables the compiler to reason about the composition and validity of transformations before any machine instructions are emitted.

### 2.3.1 The `:mereology` partition: The Language of Parts

The `:mereology` partition introduces the formal signature of parthood as the foundational structural primitive of the library. Following the mereological grounding of Stanisław Leśniewski, we treat parthood not as naive set membership, but as a primitive binary relation ( $\leq$ ) that governs the structural identity of an entity. **\*\*If the Morphism is the atom of composability, then Mereology is the atom of hierarchy\*\***, providing the skeletal vocabulary required for any downstream category to reason about sub-structures.

By utilizing the signature discovery established in the previous section, the library reified parthood as a **Morphic Constraint**. A relation  $R : A \times A \rightarrow \Omega$  is verified as a valid mereological inhabitant only if it satisfies the axioms of a **Partial Order**. As shown in Listing 9, the `IsPartRelation` concept acts as a formal verification gate, enforcing reflexivity ( $a \leq a$ ), transitivity ( $a \leq b \wedge b \leq c \implies a \leq c$ ), and antisymmetry ( $a \leq b \wedge b \leq a \implies a = b$ ).

Crucially, the library does not yet bind these axioms to specific hardware operators like `std::less_equal`. This abstraction allows the `:logic` partition (2.3.2) to address the "Reflexivity Failure" observed in IEEE 754 floating-point species, where the singularity `NaN <= NaN` evaluates to `false`. By isolating these axioms in Level 0, we ensure that the subsequent synthesis of Set Theory in Level 4 inherits a verified, algebraic definition of "part-of" relationships.

```
/**
 * @brief The Mereological Signature: An abstract Partial Order.
 * @tparam R The Relation f: A x A -> Omega (Subobject Classifier)
 */
template <typename R, typename T>
concept IsPartRelation = IsArrow<R, T, T, bool> &&
    requires(R rel, T a, T b, T c) {
    // 1. Reflexivity: Pars partis est pars totius.
    static_assert(rel(a, a), "Reflexivity_␣failed.");

    // 2. Transitivity: Hierarchical propagation.
    { rel(a, b) && rel(b, c) } -> std::convertible_to<bool>;

    // 3. Antisymmetry: Identity of indiscernible parts.
    { rel(a, b) && rel(b, a) } -> std::same_as<bool>;
};

// SANITY CHECK: Integers satisfy the Mereological Signature.
static_assert(IsPartRelation<std::less_equal<int>, int>,
    "Verified:␣Integrals␣form␣a␣valid␣Mereological␣hierarchy.");

// WARNING: Floats fail the Atomic Floor due to NaN reflexivity holes.
// static_assert(IsPartRelation<std::less_equal<double>, double>); // Fails!
```

Listing 9: The Language of Parts: Verifying the Mereological Order.

### 2.3.2 The `:logic` partition: The Subobject Classifier ( $\Omega$ )

Before a “Body” (Set) or a “Path” (Sequence) can be reified, the `dedekind` library establishes the “Rules of Presence” through the **Subobject Classifier**  $\Omega$ . Following the algebraic logic of Rasiowa [20], we treat every logical system as a specific type of abstract algebra where truth-values act as the objects of a **Topos**. By reifying these systems via the `LogicalSpecies` concept, we prevent the “leaky abstractions” inherent in machine-level `bool` types—such as implicit integral promotion—while allowing for pluggable logical universes.

In the `ClassicalLogic` species,  $\Omega$  is the binary prime  $\{True, False\}$  (see Table 3), corresponding to the standard Boolean Topos where the **Law of Excluded Middle** (LEM) holds. However, to honestly model computational undecidability or the “truth-holes” inherent in floating-point boundaries, we introduce the `TernaryLogic` species. Based on Kleene’s strong logic of indeterminacy ( $K3$ ) [19], this logic introduces an **Unknown** state (0). Within `dedekind`, these logics are formalized as **Rigs**; the *Join* ( $\vee$ ) serves as additive composition (Supremum), while the *Meet* ( $\wedge$ ) serves as multiplicative composition (Infimum).

The necessity of `TernaryLogic` is most evident when observing the breakdown of the LEM in floating-point arithmetic. Following IEEE 754 [12], the `NaN` (Not-a-Number) state represents a singularity where standard relational morphisms fail to provide a binary partition. For a predicate  $P(x) := (x \leq 1.0)$  where  $x = \text{NaN}$ , classical C++ evaluates both  $P(x)$  and its logical complement  $\neg P(x)$  as `false`. This results in the paradox  $P \vee \neg P = \perp$ , a violation of classical tautology that `dedekind` resolves by elevating the subobject classification to  $\Omega_{K3}$ .

Table 3: Truth Tables for Classical ( $\Omega_{\mathbb{B}}$ ) and Kleene ( $\Omega_{K3}$ ) Toposes.

Table 4: \*

Classical ( $\Omega = \{0, 1\}$ )

| $A$ | $B$ | $A \wedge B$ | $A \vee B$ | $\neg A$ |
|-----|-----|--------------|------------|----------|
| 0   | 0   | 0            | 0          | 1        |
| 0   | 1   | 0            | 1          | 1        |
| 1   | 0   | 0            | 1          | 0        |
| 1   | 1   | 1            | 1          | 0        |

Table 5: \*

Kleene K3 ( $\Omega = \{-1, 0, 1\}$ )

| $A$ | $B$ | $A \wedge B$ | $A \vee B$ | $\neg A$ |
|-----|-----|--------------|------------|----------|
| -1  | -1  | -1           | -1         | 1        |
| -1  | 0   | -1           | 0          | 1        |
| -1  | 1   | -1           | 1          | 1        |
| 0   | 0   | 0            | 0          | 0        |
| 0   | 1   | 0            | 1          | 0        |
| 1   | 0   | 0            | 1          | -1       |
| 1   | 1   | 1            | 1          | -1       |

**The Lipschitz Boundary.** For signed integral types, the mapping from the infinite ring of integers to finite machine representations is only stable within a defined “Lipschitz domain”  $[-2^{N-1}, 2^{N-1} - 1]$ . By utilizing compiler built-ins (e.g., `__builtin_add_overflow`), the `SubobjectClassifier` detects boundary breaches and maps them to the **Unknown** state, ensuring that downstream transformations treat these singularities as formal subobject exclusions rather than undefined behavior.

```
// In the Dedekind Ternary Topos (Omega):
Ternary P      = Classifier<double>::eval(x <= 1.0); // Unknown
Ternary NOT_P  = !P;                                // Unknown
Ternary LEM    = P || NOT_P;                        // Unknown

/** @brief Lipschitz Boundary Check for Signed Integers */
template <typename Op, typename T>
static constexpr Omega evaluate_arithmetic(T a, T b, Op op) {
    T result;
    if (__builtin_add_overflow(a, b, &result)) {
        return L::Unknown; // Lipschitz boundary breached
    }
    return L::True;
}
```

Listing 10: The NaN Truth-Hole and Lipschitz Boundary.

### 2.3.3 Posetal Categories (:posetal)

The `:posetal` partition matures the abstract parthood defined in `:mereology` into a formal **Posetal Category**. It is critical to note that while traditional mathematics defines a Poset over a *set*, the `dedekind` build graph establishes these relations at Level 0—before the `:etcs` partition (Section 2.3.1) provides the set-theoretic axioms. Consequently, we treat posetal structures not as collections of elements, but as **Skeletal Categories** where, for any two objects  $A$  and  $B$ , there exists *at most one* morphism  $f : A \rightarrow B$ .

In this skeletal framework, the existence of a unique arrow is governed by the parthood relation ( $\leq$ ) as evaluated by the chosen Subobject Classifier  $\Omega$ . By anchoring these “Skeletal Orders” early, we provide the order-theoretic foundation required for the Lattices in Level 1 and the Dedekind-completeness proofs in Level 4. The `IsPosetal` concept (Listing 11) verifies that the relation is decidable within the species’ specific logic, transforming a raw type into a structured coordinate system.

```
/**
 * @brief The Posetal Concept: A Skeletal Category.
```

```

* @tparam T The Species (the objects of the category).
* @tparam Rel The Partial Order (the unique morphisms).
*/
template <typename T, typename Rel>
concept IsPosetal = IsPartRelation<Rel, T> &&
    requires (Rel rel, T a, T b) {
        // The relation must return a value from the Species' Topos Omega
        { rel(a, b) } -> std::same_as<typename Classifier<T>::type>;
    };

// Verification: Standard integers under Classical Logic form a Posetal structure.
static_assert(IsPosetal<T, std::less_equal<int>>,
    "Verified: Integrals form a valid Posetal Category.");

// Structural Pruning: The compiler resolves categorical transitivity.
template <typename T, typename Rel>
requires IsPosetal<T, Rel>
constexpr auto check_chain(T a, T b, T c) {
    // If the arrows A->B and B->C exist...
    if constexpr (Rel{}(a, b) && Rel{}(b, c)) {
        // ...the arrow A->C is a categorical necessity.
        static_assert(Rel{}(a, c), "Categorical Transitivity Violation!");
        return true;
    }
    return false;
}

```

Listing 11: The Posetal Category: Orders as Unique Morphisms.

### 2.3.4 Small Categories (:small)

The `:small` partition formalises the transition from relational parthood to the general **Small Category**. While a *Posetal Category* (Section 2.2.5) restricts the morphism count between objects to at most one, the `IsSmallCategory` concept removes this constraint, reifying the axioms of **Categorical Composition** and **Identity**. In the `dedekind` ontology, a small category is treated as a finite registry of types and arrows that are visible to the compiler's proof-search engine.

By anchoring these axioms in Level 0, we provide the skeletal structure required for the *Functors* in Level 2. Crucially, the library enforces that for every object  $X$  in the category, there exists an identity morphism  $id_X$  and that the composition of arrows  $f : A \rightarrow B$  and  $g : B \rightarrow C$  is associative. As shown in Listing 12, this allows the compiler to verify that a transformation pipeline is not merely a sequence of function calls, but a mathematically valid path through the species' graph.

```

/**
* @brief The Small Category Concept: Composition and Identity laws.
* @tparam Cat The Category registry.
*/
template <typename Cat>
concept IsSmallCategory = requires {
    // 1. Identity: Every object X must have an identity arrow id_X.
    requires requires (typename Cat::Object X) {
        { Cat::identity(X) } -> IsArrow<typename Cat::Object, typename Cat::Object>;
    };

    // 2. Composition: f: A -> B and g: B -> C must compose to h: A -> C.
    requires requires (typename Cat::Arrow f, typename Cat::Arrow g) {

```

```

        requires std::same_as<typename SpeciesTraits<f>::Codomain,
                               typename SpeciesTraits<g>::Domain>;
        { Cat::compose(g, f) } -> IsArrow;
    };
};

// SANITY CHECK: The 'Atomic' category of C++ types satisfies Smallness.
static_assert(IsSmallCategory<CPlusPlusTypeCategory>);

```

Listing 12: The Small Category: Composition and Identity.

### 2.3.5 Universal Products and Coproducts (:cartesian)

To support multi-variable transformations and formal choice, the `:cartesian` partition reifies `std::pair` and `std::variant` through their universal properties. In the `dedekind` topos, a **Product** ( $A \times B$ ) is not merely a passive data container, but a structural requirement for the existence of unique projections  $\pi_1$  and  $\pi_2$ . By mapping the categorical product to `std::pair`, we establish the formal basis for binary algebraic operations—modeled as morphisms of the form  $X \times X \rightarrow X$ —ensuring that the “Rules of Both” are verified as categorical invariants.

Dually, the **Coproduct** ( $A + B$ ) is reified through `std::variant`, representing the “Rules of Choice.” This allows for the categorical modeling of branching logic: a morphism out of a coproduct  $f : A + B \rightarrow C$  is uniquely determined by a pair of morphisms ( $f_A : A \rightarrow C, f_B : B \rightarrow C$ ) via the mediator. By anchoring these universal constructions in the standard library’s vocabulary, the library enables the `:algebra_total` partition to define complex signatures while maintaining the zero-overhead performance of C++ machine types.

```

/**
 * @brief Categorical Product Verification:
 * Ensures T satisfies universal projection properties for A and B.
 */
template <typename T, typename A, typename B>
concept IsProduct = requires(T product) {
    { project<0>(product) } -> std::same_as<A>;
    { project<1>(product) } -> std::same_as<B>;
};

// Verification of the Categorical Product (std::pair)
static_assert(IsProduct<std::pair<int, int>, int, int>,
              "Verified: std::pair is a valid Categorical Product.");

/**
 * @brief Categorical Coproduct Verification:
 * Ensures T satisfies the Rule of Choice (Injectors).
 */
template <typename T, typename A, typename B>
concept IsCoproduct = requires(A a, B b) {
    { inject<A>(a) } -> std::same_as<T>;
    { inject<B>(b) } -> std::same_as<T>;
};

// Verification of the Categorical Coproduct (std::variant)
static_assert(IsCoproduct<std::variant<int, bool>, int, bool>,
              "Verified: std::variant is a valid Categorical Coproduct.");

// Binary Morphism Signature: X * X -> X
using Addition = Morphism<std::pair<int, int>, int, std::plus<int>>;

```

```
| static_assert(IsArrow<Addition, std::pair<int, int>, int>);
```

Listing 13: The Universal Bridge: Mapping Products and Coproducts to C++ Primitives.

### 2.3.6 Boundary Objects (:limit)

The `:limit` partition defines the logical boundaries of the category: the **Initial** (0) and **Terminal** (1) objects. Following the structuralist approach advocated by Lawvere [21], these are defined not by their internal contents, but by the existence of unique morphisms to or from these boundaries. In the C++23 context, we reify the Terminal object as `std::monostate` (or a unit type), representing a state of pure existence without information.

This abstraction allows for the formal definition of “elements” as morphisms from the terminal object ( $1 \rightarrow X$ ). In `dedekind`, a constant value is not a static literal, but a *global element*—a mapping that selects a specific inhabitant of a species. Dually, the Initial object (0) represents the absolute annihilator; a morphism from the initial object  $0 \rightarrow X$  is the categorical equivalent of an unreachable code path. By anchoring these boundaries early in the build graph, we provide the formal basis for the `:etcs` partition to reason about set membership as a structural relation.

```
/** @brief The Terminal Object (1): Pure existence. */
using One = std::monostate;
static_assert(IsTerminalObject<One>);

/** @brief Defining a "Global Element" as a Morphism 1 -> X */
auto five = Morphism<One, int>([](One) { return 5; });
static_assert(IsArrow<decltype(five), One, int>);

/** @brief Initial Object (0): The unreachable annihilator. */
using Zero = std::nullptr_t; // Categorical void / Unreachable
static_assert(IsInitialObject<Zero>);

// Verification: Every species T has a unique morphism T -> 1 (Terminality)
static_assert(HasUniqueMorphismTo<int, One>);
```

Listing 14: The Boundary Objects: Reifying 0 and 1 in C++23.

### 2.3.7 Universal Constructions and Pullbacks (:pullback)

To support relational reasoning and intersections, the `:pullback` partition introduces the universal property of the `**Pullback**`. As noted in Pierce [**pierce1991basic**], pullbacks are the categorical tool for handling constraints and fiber products. By reifying the pullback, the library enables the formal definition of **Equalizers**, which are essential for identifying the rounding-error singularities in the `:numeric` partition.

In the `dedekind` ontology, the pullback  $P$  of two morphisms  $f : A \rightarrow C$  and  $g : B \rightarrow C$  represents the subset of the product  $A \times B$  where the morphisms coincide ( $f(a) = g(b)$ ). This construction is the categorical engine for set intersection; by hoisting this requirement into the type-system, we allow the compiler to perform `**Symbolic Equalization**`, identifying when two computational paths are structurally identical or divergent before any floating-point operation is executed.

```
/** @brief Pullback P of f: A -> C and g: B -> C */
template <typename f, typename g>
using Pullback = Set<Product<Domain<f>, Domain<g>>, [](auto pair) {
    return f{}(project<0>(pair)) == g{}(project<1>(pair));
}>;

// An Equalizer is a pullback of f and g over the same codomain
```



```
using Eq = Equalizer<f, g>;
static_assert(IsLimit<Eq>);
```

Listing 15: Pullbacks and Equalizers: The Engine of Constraint.

## 2.4 Level 1: Space and Action (dedekind.category)

The transition from Level 0 to Level 1 represents the "Maturation of Species." Following the Polish School of Logic as articulated by Ajdukiewicz [22], we treat mathematical structures as active apparatuses—or conceptual orientations—that determine the image of the computational world. This aligns with the restriction category approach [23], where totality is not a default state but a property verified against a specific subobject support.

### 2.4.1 The :total partition: The Property of Totality

The :total partition formalizes *Morphic Totality*. In our ontology, a morphism  $f : A \rightarrow B$  is verified as a `TotalMorphism` only if it is defined across its entire domain without triggering a subobject exclusion (e.g., overflow or division by zero). Following the restriction category framework [23], we treat totality as a property verified against a specific subobject support, ensuring that the “active apparatus” of the morphism is physically stable within the machine’s constraints.

```
/** @brief Formal verification that a morphism is defined for the entire species. */
export template <typename F, typename A, typename B>
concept IsTotalArrow = IsArrow<F, A, B> && requires(F f, A x) {
    // Totality check: The SubobjectClassifier must return 'True'
    requires std::same_as<typename SubobjectClassifier<A>::Omega, bool>;
    { SubobjectClassifier<A>::evaluate_total(x) } -> std::same_as<bool>;
};

export template <IsSpecies A, IsSpecies B, typename Func>
struct TotalMorphism : Morphism<A, B, Func> {
    static_assert(IsTotalArrow<Func, A, B>, "Totality_Error:_Function_is_not_total.");
};
```

Listing 16: Totality Verification and the Total Morphism wrapper.

### 2.4.2 The :algebra\_total partition: The Laws of Harmony

Building upon the total mappings established in Level 1a, the :algebra\_total partition materializes the “Pure Ideals” of abstract algebra. In this partition, species “mature” as they gain axioms, moving from the primal **Magma** to the perfect symmetry of an **Abelian Group**. A key architectural decision in *dedekind* is the decoupling of *associativity* from *invertibility*: while a **Monoid** provides an identity to an associative Semigroup, a **Unital Magma** (or **Loop** if invertible) provides an identity even when associativity is rejected [24]. A **Group** is thus reified as the perfect synthesis of these two paths, as shown in Table 6.

This leads to the **Honest Rejection** policy: while `unsigned int` addition is verified as a total Abelian Group ( $\mathbb{Z}/2^n\mathbb{Z}, +$ ), `signed int` is relegated to a **Unital Magma**. The library discovers the identity element (0) as a safe invariant, yet rejects the **Associative** and **Distributive** axioms. Because signed overflow constitutes undefined behavior, the C++ abstract machine cannot guarantee structural harmony under re-ordering; the library correctly identifies such operations as **Hazards** [25], preserving the integrity of the total domain.

```
/** @section Truth_Anchors */

// 1. Verification of the Modular Ideal ( $\mathbb{Z}/2^n\mathbb{Z}, +$ )
```

Table 6: The Path to Symmetry: Algebraic Species in `dedekind.category`. The **Group** represents the synthesis of the associative path (Monoid) and the invertible path (Loop).

| Concept              | Structure                     | Axiomatic Requirements                          | C++ Concept                 |
|----------------------|-------------------------------|-------------------------------------------------|-----------------------------|
| <b>Magma</b>         | $\langle T, \circ \rangle$    | Closure: $\forall a, b \in T : a \circ b \in T$ | <code>IsMagma</code>        |
| <b>Semigroup</b>     | $\langle T, \circ \rangle$    | Magma + Associativity                           | <code>IsSemigroup</code>    |
| <b>Monoid</b>        | $\langle T, \circ, e \rangle$ | Semigroup + Identity                            | <code>IsMonoid</code>       |
| <b>Loop</b>          | $\langle T, \circ, e \rangle$ | Magma + Identity + Unique Division              | <code>IsLoop</code>         |
| <b>Group</b>         | $\langle T, \circ, e \rangle$ | Monoid + Loop                                   | <code>IsGroup</code>        |
| <b>Abelian Group</b> | $\langle T, \circ, e \rangle$ | Group + Commutativity                           | <code>IsAbelianGroup</code> |

```
static_assert(IsAbelianGroup<unsigned int, std::plus<>>,
    "Axiom_Error: Unsigned addition is a perfect Abelian Group.");

// 2. The Honest Rejection (Hazardous Domain)
// Signed addition provides an identity (Unital Magma),
// but is rejected as a Semigroup due to the UB overflow hazard.
static_assert(has_identity_v<int, std::plus<>>,
    "Maturation: Signed identity(0) is safe and discoverable.");

static_assert(!is_associative_v<int, std::plus<>>,
    "Honesty_Check: Signed associativity rejected due to UB.");

static_assert(!IsSemigroup<int, std::plus<>>,
    "Taxonomy: Signed int addition is not a Semigroup.");

// 3. The Boolean Ring (Characteristic 2)
static_assert(characteristic_v<bool> == 2,
    "Axiom_Error: Boolean Ring must have characteristic 2.");

static_assert(is_idempotent_v<bool, std::logical_or<>>,
    "Axiom_Error: Boolean OR must be idempotent.");

// 4. Distributive Handshake
static_assert(is_distributive_v<unsigned int, std::multiplies<>, std::plus<>>,
    "Axiom_Error: Unsigned * must distribute over +.");

static_assert(!is_distributive_v<int, std::multiplies<>, std::plus<>>,
    "Honesty_Check: Distributivity rejected for signed hazards.");
```

Listing 17: Structural Maturation and the Honest Rejection in `dedekind.category`.

### 2.4.3 Monoid Actions and Modules (:actions)

The `:actions` partition materializes the concept of a Species being acted upon by an Algebraic Structure. This partition defines the `IsSemiModule` and `IsModule` concepts, where a **Monoid** (from `:algebra_total`) defines a transformation rule over a target Species.

Following the pedagogical tradition of Stambach [26], we treat linear algebra not as a series of matrix operations, but as the study of structured actions. In this ontology, we distinguish between **Internal Composition** (operations within a species) and **External Influence** (actions upon a species). This "coordinate-free" approach [27] is critical for transitioning from pure algebraic laws to the modeling of transformations across spaces, establishing the formal basis for the linear transformations and tensor structures later introduced in the `:geometry` module.

#### 2.4.4 The `:functor` partition: Mappings between Categories

Following the fractal progression of the library, the `:functor` partition reifies the algebraic laws of Level 1 to act upon the categories themselves. In `dedekind`, a `Functor` is treated not as a container, but as a structural morphism  $F : \mathcal{C} \rightarrow \mathcal{D}$  that preserves the composition of arrows and the identity of objects. By elevating the notion of mapping to a category-level operation, we allow the library to reason about transformations across different ontological domains.

The implementation centers on a high-level `fmap` dispatcher. Leveraging C++23 template-template parameters (e.g., `template <typename> typename F`), the dispatcher ensures that functorial mapping is a zero-overhead operation. This architecture shields the user from the underlying machine representation while maintaining strict categorical rigor. By verifying that  $F(g \circ f) = F(g) \circ F(f)$  at compile-time, the compiler ensures that the structural integrity of the mathematical model remains invariant under transformation.

```
/**
 * @brief The Functorial Dispatcher (fmap).
 * Maps an Arrow (f: A -> B) to a lifted Arrow (F(f): F<A> -> F<B>).
 */
export template <template <typename> typename F, IsArrow Arrow>
constexpr auto fmap(Arrow f) {
    return lifted_morphism<F>(f);
}

// Verification: Functors must preserve the Identity Morphism.
// Here, 'F' represents a generic skeletal functor.
static_assert(PreservesIdentity<F, int>);
```

Listing 18: Functorial Mapping: Mappings as zero-overhead dispatchers.

#### 2.4.5 The `:natural` partition: Natural Transformations

To complete the Level 2 bridge, the `:natural` partition introduces **Natural Transformations**—morphisms between functors. This provides the mathematical basis for “component-wise” transformations that remain invariant under functorial mapping, satisfying the naturality square:  $\alpha_Y \circ F(f) = G(f) \circ \alpha_X$ . By formalizing these mappings, the library ensures that structural changes—such as the transition from a skeletal species to an effectful one—are consistent across all objects in the category.

The `:natural` partition is a strict technical prerequisite for the `:monad` partition, as it establishes the formal requirements for the unit ( $\eta$ ) and multiplication ( $\mu$ ) transformations. In `dedekind`, naturality is not merely a convention but a verified invariant. Leveraging C++23 Concepts, we implement the `IsNaturalTransformation` constraint to verify that a structural mapping  $\alpha$  commutes with the `fmap` of its source and target functors. This verification ensures that the underlying structural integrity is maintained across different functorial contexts without “species-leak” or logical drift.

```
/**
 * @concept IsNaturalTransformation
 * @brief Formal verification of the naturality square.
 */
export template <typename Alpha, template <typename> typename F,
                template <typename> typename G>
concept IsNaturalTransformation = requires(Alpha alpha) {
    // For any morphism f: A -> B, the following must commute:
    // alpha<B> o fmap<F>(f) == fmap<G>(f) o alpha<A>
    requires CommutesWithFmap<Alpha, F, G>;
};
```

```

// Skeletal unit transformation (eta): Id -> F
export template <template <typename> typename F>
struct UnitTransformation;

```

Listing 19: The Naturality Square: Verifying consistency between Functors.

#### 2.4.6 The :monad partition: The $\Omega$ -Monad

While Level 1 defines absolute laws, real-world computation requires a mechanism for handling “effectful” transformations. Following the constructions of Kleisli [28] and Moggi [6], we implement the **Kleisli Triple** ( $\eta, \gg=$ ) as the fundamental unit of composition. In the :monad partition, we unify the “Rules of Presence” from the :logic topos with this machinery to define the  $\Omega$ -Monad.

As illustrated in Listing 20, we extract the `IsPotential` concept to create a structural isomorphism between standard C++ types and the `dedekind Partial<T>` container. To resolve the lack of native higher-kinded types (HKTs) in C++, we utilize an “Action-First” bootstrapping strategy: the `fmap` functorial mapping is derived as an *epi-phenomenon* of the underlying monadic “Push” ( $m \gg= \eta \circ f$ ). This ensures that any species providing a `bind` operator ( $\gg=$ ) automatically matures into a *Functor* through *Synthetic Inference*.

```

export template <typename R>
concept IsPotential = requires(R r) {
    typename R::value_type;
    { *r } -> std::convertible_to<typename R::value_type&>;
    { presence_of(r) } -> LogicalValue; // Bridge to Omega
};

// The Universal Highway: Works for any IsPotential result
export template <IsPotential M, typename F>
constexpr auto operator>>=(const M& m, F f) {
    using L = typename GetLogic<M>::type;
    if (presence_of(m) != L::True) return empty<M>(presence_of(m));
    return f(*m);
}

```

Listing 20: The  $\Omega$ -Monad: Structural Parity and the Universal Highway.

#### 2.4.7 The :kleisli partition: The Universal Highway

The :kleisli partition introduces the formal machinery for *Highway composition* (operator  $\gg$ ). Following the functional tradition of the “Fish Operator” ( $\gg=$ ) [29], `dedekind` adapts this to the C++23 context by leveraging the left-associativity of `operator>>` to model the directional vector of computation across the ontology.

To resolve the lack of native higher-kinded types (HKTs) in C++, the library utilizes an *Action-First* bootstrapping strategy. By defining Monads as Extension Systems (Kleisli Triples consisting of  $\eta$  and  $\gg=$ ), we utilize operator overloading on concrete objects to trigger Argument-Dependent Lookup (ADL) at the call site. This allows the Level 0 foundation to instantiate a derived `fmap` as an *epi-phenomenon* of the underlying monadic “Push” ( $m \gg= \eta \circ f$ ).

The resulting *Highway Notation* reifies categorical data flows as directional vectors:

- `value >> into<F>` : Represents  $\eta$  (Unit), lifting a species into a context.
- `box << extract<F>` : Represents  $\varepsilon$  (Counit), sampling a species from a context.
- `box >>= f` : Represents Bind, the forward monadic composition.

- `box <= f` : Represents Extend, the contextual comonadic composition.

Within the verified module graph, the LLVM backend utilizes this structural transparency to perform *Highway Pruning*. If a total function is composed with a partial operation statically known to return an **Unknown** or NaN state, the compiler identifies the unreachable path and optimizes the Highway into a branch-free instruction sequence.

#### 2.4.8 The `:partial` partition: The Pragmatic Rescue

The `:partial` partition performs the “Categorical Rescue” of species that failed the strict requirements of the total domain. This stage formalises the transition from the *Platonic* to the *Pragmatic*: acknowledging that while hardware-level operations may be “broken” (e.g., via overflow or division by zero), they are **consistently broken**. Following the formalisations of the *nLab* community [30], we treat these consistent failures not as exceptions, but as structured data.

Following topos-theoretic foundations [19], we treat a partial morphism  $f : A \multimap B$  as a total mapping from a *subobject*  $S \hookrightarrow A$ . Here,  $S$  represents the **Support** defined by the characteristic function  $\chi : A \rightarrow \Omega_{K3}$  provided by the `:logic` partition. Within this restricted domain, we define *locally total* structures, such as **Partial Groups** or **Partial Monoids**. These structures guarantee that the “Laws of Harmony” hold strictly on the support  $U = \{a \in A \mid \chi(a) = \top\}$ , ensuring that even when navigating the messy reality of hardware, we remain within a rigorous algebraic landscape.

By reifying the **Partial**<T> “Rescue Pod” as an inhabitant of the  $\Omega$ -Monad, the library enables the composition of these partial morphisms via the universal Highway established in the `:kleisli` partition. The bind operator (`>>=`) acts as the monadic gatekeeper: it consumes the **Ternary** presence, ensuring that any “Truth-Hole” is propagated through the Kleisli chain. This transforms inconsistent machine behavior into a rigorous, composable algebraic flow.

```
// The "Rescue": int addition is recognized as a Partial Abelian Group.
static_assert(IsPartialAbelianGroup<int, std::plus<int>>);

// Composition in the Kleisli Category via the Highway:
auto f = partial_plus(1, 2);           // Result: 3 (True)
auto g = partial_plus(f, INT_MAX);    // Result: Unknown (Overflow)
auto h = partial_plus(g, 1);           // Result: Unknown (Propagation)
```

Listing 21: The Kleisli Lift: Reclaiming the Lipschitz Boundary.

#### 2.4.9 The `:algebra_partial` partition: The Generalised Reality

Following the monadic Kleisli lift, the `:algebra_partial` partition generalizes the algebraic laws of Level 1 to handle the “broken” reality of hardware. In this partition, the strict requirements of absolute closure and identity are relaxed into **Partial Monoids** and **Partial Groups**. Unlike their total counterparts, these structures are only guaranteed to hold within the specific support  $S \hookrightarrow A$  provided by the `:logic` and `:partial` partitions.

The core objective of this partition is to verify that a species is “consistently broken.” By reifying partiality through the  $\Omega$ -Monad, **dedekind** ensures that if a result is defined (i.e., its presence is  $\top$ ), it must still satisfy the required algebraic identities. For instance, a **Partial Monoid** must remain associative wherever its composition is defined. This allows the library’s dispatcher to safely propagate **Unknown** or NaN states through the algebraic pipeline without triggering undefined behavior. By enforcing these laws within the support, we ensure that the “Laws of Harmony” are not discarded, but are instead localized, allowing the compiler to perform structural pruning even in the presence of indeterminate data.

```

// Verification: A Partial Monoid is associative on its support.
export template <typename T, typename Op>
concept IsPartialMonoid = IsPartialMagma<T, Op> &&
                        HasLocalIdentity<T, Op> &&
                        IsLocallyAssociative<T, Op>;

// Signed integers: Total Magma (REJECTED), Partial Abelian Group (SUCCESS).
static_assert(IsPartialAbelianGroup<int, std::plus<int>>);

```

Listing 22: Local Totality: Verifying the Laws of Harmony on the Support.

#### 2.4.10 The :numeric partition: The Material Landing

The :numeric partition represents the final “landing” of the abstract theory back onto the hardware substrate. While the :partial and :algebra\_partial partitions establish the general laws for inconsistent structures, the :numeric partition provides the specific mapping for IEEE 754 floating-point types and bounded integers.

In this partition, the library utilizes the `Partial<T>` rescue pod to formally model the non-associative nature of `double` and the non-reflexive behavior of `NaN`. By anchoring these hardware-specific behaviors as first-class inhabitants of a partial monoid, we ensure that the “messy reality” of rounding errors and overflows is not an exception to the theory, but a verified and traceable component of the categorical flow. This is achieved through “Safety Certificates” for specific machine species, which allow the `dedekind` dispatcher to dynamically enable or disable algebraic optimizations (such as expression reordering) based on the material reality of the underlying type.

Finally, we encounter the curious case of *Numerical Species* (IEEE 754). These types are effectively categories where the subobject classifier  $\Omega$  is internalized within the species itself via states such as `NaN` and `Inf`. While these structures appear to be total morphisms  $\mu : \mathbb{R}_f \times \mathbb{R}_f \rightarrow \mathbb{R}_f$  in their machine signature, they are semantically partial. Specifically, they fail the **Equalizer** conditions required for bit-perfect associativity; the rounding error ensures that  $(a + b) + c$  and  $a + (b + c)$  are not generally isomorphic. This progression—from the absolute (*Total*), through the filtered (*Partial*), to the exceptional (*Numerical*)—allows us to build a rigorous ontology that respects both categorical rigor and the physical constraints of the C++ host language.

```

// IEEE 754 Addition: Partial Abelian Group (SUCCESS).
static_assert(IsPartialAbelianGroup<double, std::plus<double>>);

// However, it is NOT a Total Monoid due to rounding error.
static_assert(!IsTotalMonoid<double, std::plus<double>>);

// Safety Certificate: Disables reassociation for floating-point.
static_assert(!is_associative_v<double, std::plus<double>>);

```

Listing 23: The Numeric Landing: Verifying the IEEE 754 Partial Monoid.

#### 2.4.11 The :etcs partition: The Set-Theoretic Interface

The final stage of the foundational build anchors Lawvere’s Elementary Theory of the Category of Sets (ETCS) [lawvere1964]. This partition defines a set not by its internal elements, but by its universal properties, utilizing C++23 Concepts to enforce well-pointedness and the existence of power objects.

The core innovation of this partition is the reification of the *Set-Builder* notation. By representing a set  $S = \{x \in X \mid P(x)\}$  as a **Stateless Constraint Type**, we transform the C++ translation unit into a symbolic laboratory.

```

// Define a predicate: x is even and greater than 10
auto P = [](auto x) { return (x % 2 == 0) && (x > 10); };

// Materialize the set at compile-time as a structural type
using EvenGreaterTen = dedekind::set<int, P>;

static_assert(dedekind::contains<EvenGreaterTen>(12));

```

Listing 24: Reifying the Set-Builder via Stateless Constraints.

By hoisting these mereological invariants into the type signature, we enable *Symbolic Equalization*. The compiler can decompose and optimize a set based on its logical essence rather than its storage layout, allowing contradictory set-builder predicates to collapse into zero-instruction machine code. This ensures that the performance of the “Naked Loop” is always mathematically justified by the prior formal certification of the ETCS axioms.

## 2.5 sets Module

For the set operations, we may also want to look at: "Monadic Intersection Types, Relationally, and Ordered <https://dl.acm.org/doi/10.1145/3777483> ... and references cited therein.

### 2.5.1 The :set partition

To ground this theoretical approach, Table 1 maps the fundamental axioms of the Elementary Theory of the Category of Sets (ETCS) – as codified in the pedagogical treatment by Lawvere and Rosebrugh [31] – to their structural equivalents in `dedekind`. This mapping ensures that the library is not merely a numerical simulation, but a direct implementation of algebraic laws within the host language’s type system. In a sense, `dedekind` hopes to be a *\*compiler\** for mathematical concepts.

Table 7: Complete Mapping of ETCS Axioms to C++23 Structural Invariants

| ETCS Axiom           | C++ Language Feature                     | Structural Invariant                      |
|----------------------|------------------------------------------|-------------------------------------------|
| 1. Composition       | <code>std::invoke / operator»</code>     | <code>IsSemigroupoid&lt;T, Op&gt;</code>  |
| 2. Identity          | <code>identity_trait&lt;T, Op&gt;</code> | <code>IsSmallCategory&lt;T, Op&gt;</code> |
| 3. Terminal Object   | <code>s(x) -&gt; Logic::top</code>       | <code>IsTerminalObject&lt;T&gt;</code>    |
| 4. Well-Pointedness  | <code>operator()</code> (Element access) | <code>IsWellPointed</code>                |
| 5. Cartesian Product | <code>std::tuple / std::pair</code>      | <code>HasProduct&lt;A, B&gt;</code>       |
| 6. Exponentiation    | <code>Lambda NTTP / std::function</code> | <code>IsInternalHom&lt;A, B&gt;</code>    |
| 7. Equalizers        | <code>dedekind::set&lt;T, P&gt;</code>   | <code>IsEqualizer&lt;f, g&gt;</code>      |
| 8. Empty Set         | <code>dedekind::EmptySet</code>          | <code>IsInitialObject&lt;T&gt;</code>     |
| 9. Infinity (NNO)    | <code>dedekind::Naturals</code>          | <code>HasNaturalNumbersObject</code>      |
| 10. Axiom of Choice  | <code>Lattice Dispatcher</code>          | <code>AxiomOfChoice</code>                |

By adopting a categorical foundation, the library provides a unified interface for both *extensional* and *intensional* representations. For finite collections, such as the `Booleans`, the system performs *extensional* materialization to compute exact results at compile-time. Conversely, for infinite or transcendental structures like the `Complex Numbers`, the library utilizes *intensional* symbolic expressions to maintain structural integrity. By embedding these categorical definitions in the C++ host language, the programmer—and crucially the *compiler*—can examine the algebraic properties of a field even when the underlying elements cannot be fully materialized as discrete machine data.

### 2.5.2 Intensional Sets and Computational Undecidability

The utility of the categorical approach is most evident when defining sets whose membership conditions are conceptually clear but computationally undecidable. Consider the set of *Normal Numbers*  $\mathcal{N}$ , defined as those real numbers whose infinite digit expansion follows a uniform distribution. Formally, a real number  $x$  is normal in base  $b$  if for every finite sequence of digits  $s$ , the limiting frequency of  $s$  appearing in the expansion of  $x$  satisfies:

$$\lim_{n \rightarrow \infty} \frac{N(s, x, n)}{n} = \frac{1}{b^{|s|}} \quad (1)$$

where  $N(s, x, n)$  denotes the number of occurrences of the string  $s$  in the first  $n$  digits of  $x$ . While we can easily define this property as a symbolic predicate within the **dedekind** ontology, the membership of specific constants like  $\pi$  or  $e$  remains a famous open problem in mathematics.

In a traditional imperative framework, such a set is unusable because membership cannot be *computed* to a boolean result. However, by treating  $\mathcal{N}$  as an *intensional* set, the **dedekind** library allows the programmer to define morphisms that take  $\mathcal{N}$  as their formal domain. This enables the **C++** compiler to verify the structural validity of a function pipeline based solely on the *definition* of normality, even while the empirical membership of  $\pi$  remains unresolved. By hoisting the predicate into a **Stateless Constraint Type**, we decouple the logical requirement of the set from the necessity of materializing its elements. This allows the compiler to treat  $\mathcal{N}$  as a valid algebraic object, facilitating symbolic reasoning about digit distribution without requiring a prior numerical proof of normality, thereby avoiding the runtime overhead of high-precision digit evaluation.

This allows the compiler to treat  $\mathcal{N}$  as a valid algebraic object, facilitating symbolic reasoning about digit distribution without requiring a prior numerical proof of normality, thereby avoiding the runtime overhead of high-precision digit evaluation.

This intensional logic extends to non-elementary functions where the antiderivative cannot be expressed in a finite combination of elementary terms. For example, the Gaussian integral  $e^{-t^2}$  defines the *Error Function* (erf) as a structural rule for area rather than a simple algebraic formula:

$$\text{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt \quad (2)$$

While traditional backends rely on lossy polynomial approximations to evaluate this integral, the **dedekind** ontology reifies  $\text{erf}(x)$  as a set of symbolic properties. This allows the **C++** compiler to resolve identities—such as the symmetry  $\text{erf}(-x) = -\text{erf}(x)$ —as zero-overhead type transformations, delaying numerical approximation until the final realization of the continuum.

### 2.5.3 Categorification of Sequences and the Continuum

While ETCS provides the axioms for discrete collections, the representation of sequences and the analytic continuum requires a transition from set-theoretic membership to functorial structures. In the **dedekind** library, we avoid the imperative indexing of terms by treating sequences as *Combinatorial Species* [8]. By defining a sequence as a functor from the category of finite sets to itself, we can express growth rules—such as the Fibonacci recurrence—as algebraic identities between functors. This allows the compiler to resolve the structural identity of a series through symbolic substitution before any machine-level summation occurs [9]. To reach the exact real arithmetic promised in Section 1, we ground our implementation of convergence in LAWVERE’s theory of enriched categories [32]. By treating metric spaces as categories enriched over the poset of non-negative reals  $[0, \infty)$ , the library reifies Cauchy sequences not as floating-point approximations, but as formal Cauchy completions [33]. This methodology ensures that the *Dedekind cut*—defining real numbers  $\mathbb{R}$  as partitions of the rationals  $\mathbb{Q}$  [31]—is not merely a runtime filter, but a verified limit of a rational sequence. This allows the LLVM backend to



treat transcendental constants like  $e$  and  $\pi$  as symbolic invariants rather than opaque numeric variables.

The library addresses the representation of transcendental constants by distinguishing between their algebraic *definition* and their numerical *realization*. While the machine may be unable to *compute* the final digit of  $\pi$ , the `dedekind` ontology can lazily *define* it as a formal limit of a rational sequence. By encoding the LEIBNIZ series for  $\pi$  or the TAYLOR expansion for  $e$ , the library treats these constants as symbolic *intensional* types.

$$e = \sum_{n=0}^{\infty} \frac{1}{n!} = 1 + 1 + \frac{1}{2} + \frac{1}{6} + \dots \quad (3)$$

This approach leverages *lazy evaluation* at the type-level, where the C++ compiler performs symbolic substitutions of the series terms before any floating-point approximation is required. By teaching the compiler the *rule* of the sequence rather than a static value, the library ensures that mathematical identities—such as  $e^0 = 1$ —can be resolved as zero-instruction compile-time proofs.

Table 8: Categorification of Sequences and Series in the Dedekind Ontology

| Mathematical Concept                          | Set-Theoretic Form                | Categorical Framework          | C++ Structural     |
|-----------------------------------------------|-----------------------------------|--------------------------------|--------------------|
| <b>Recursive Sequence</b>                     | $F_n = F_{n-1} + F_{n-2}$         | Combinatorial Species [8]      | Functorial Coprodu |
| <b>Formal Power Series</b>                    | $\sum a_n x^n$                    | Analytic Functors [9]          | Polynomial Species |
| <b>Metric/Distance</b>                        | $ x - y  < \epsilon$              | Lawvere Enriched Category [32] | Enriched Hom-sets  |
| <b>Cauchy Convergence</b>                     | $\lim_{n \rightarrow \infty} a_n$ | Cauchy Completion [33]         | Limit / Terminal O |
| <b>Real Numbers (<math>\mathbb{R}</math>)</b> | Dedekind Cuts / Cauchy            | Archimedian Completion [19]    | Adjoint Functors   |

xxxxxx

While ETCS provides the axioms for discrete collections, the representation of sequences and the analytic continuum requires a transition from set-theoretic membership to functorial structures. In the `dedekind` library, we avoid the imperative indexing of terms by treating sequences as Combinatorial Species [8]. By defining a sequence as a functor from the category of finite sets to itself, we can express growth rules—such as the Fibonacci recurrence—as algebraic identities between functors. This allows the compiler to resolve the "combinatorial essence" of a series through symbolic substitution before any machine-level summation occurs [9]. To reach the exact real arithmetic promised in Section 1, we ground our implementation of convergence in Lawvere's theory of enriched categories [32]. By treating metric spaces as categories enriched over the poset of non-negative reals  $[0, \infty)$ , the library reifies Cauchy sequences not as floating-point approximations, but as formal Cauchy completions [33]. This methodology ensures that the Dedekind cut – defining real numbers as partitions of the rationals [31] – is not merely a runtime filter, but a verified limit of a rational sequence, allowing the LLVM backend to treat transcendental constants like  $e$  and  $\pi$  as first-class structural invariants.

The library addresses the representation of transcendental constants by distinguishing between their algebraic *definition* and their numerical *realization*. While the machine may be unable to *compute* the final digit of  $\pi$ , the `dedekind` ontology can lazily *define* it as a formal limit of a rational sequence. By encoding the Leibniz series for  $\pi$  or the Taylor expansion for  $e$ , the library treats these constants as symbolic *intensional* types.

$$e = \sum_{n=0}^{\infty} \frac{1}{n!} = 1 + 1 + \frac{1}{2} + \frac{1}{6} + \dots \quad (4)$$

This approach leverages *lazy evaluation* at the type-level, where the C++ compiler performs symbolic substitutions of the series terms before any floating-point approximation is required.

By teaching the compiler the *rule* of the sequence rather than a static value, the library ensures that mathematical identities—such as  $e^0 = 1$ —can be resolved as zero-instruction compile-time proofs.

Table 9: Categorification of Sequences and Series in the Dedekind Ontology

| Mathematical Concept                          | Set-Theoretic Form                | Categorical Framework          | C++ Structural          |
|-----------------------------------------------|-----------------------------------|--------------------------------|-------------------------|
| <b>Recursive Sequence</b>                     | $F_n = F_{n-1} + F_{n-2}$         | Combinatorial Species [8]      | Functorial Coproduct    |
| <b>Formal Power Series</b>                    | $\sum a_n x^n$                    | Analytic Functors [9]          | Polynomial Species      |
| <b>Metric/Distance</b>                        | $ x - y  < \epsilon$              | Lawvere Enriched Category [32] | Enriched Hom-sets       |
| <b>Cauchy Convergence</b>                     | $\lim_{n \rightarrow \infty} a_n$ | Cauchy Completion [33]         | Limit / Terminal Object |
| <b>Real Numbers (<math>\mathbb{R}</math>)</b> | Dedekind Cuts / Cauchy            | Archimedean Completion [19]    | Adjoint Functors        |

This mapping ensures that our software implementation is not merely a *simulation* of set theory, but a direct *instantiation* of its algebraic laws.

The structural mereology [34] approach in **dedekind** defines sets via formal relations—union ( $X \cup Y$ ), intersection ( $X \cap Y$ ), and parthood ( $X \subseteq Y$ ) and last but not least element ( $x \in Y$ )—prior to implementation. This mirrors structuralist foundations such as Lawvere’s Elementary Theory of the Category of Sets (ETCS) [35] and Cartesian Closed Categories (CCCs) [36, 37]. By prioritizing *structural transparency*, this methodology fundamentally departs from object-oriented encapsulation. In a sense, it encodes the well known rule of thumb of "composition over inheritance" [38]. In a standard `std::set`, the internal representation is an opaque implementation detail, invisible to the type system’s logic. Conversely, in **dedekind**, the *type is the definition*. Because the set’s constituents—predicates, universes, and mereological relations—are hoisted into the type signature as first-class members, the C++ type system and the **CMake** toolchain function as a Constraint Satisfaction Engine (CSE). This allows the compiler to inspect, decompose, and optimize a set based on its logical essence rather than its storage layout, trading the 'black-box' safety of encapsulation for the 'white-box' provability of formal mereology.

```
#include <concepts>
#include <type_traits>

// Verifies the "Element Of" relation (x in Y)
template <typename T, typename S>
concept IsElementOf = requires(T x, S s) {
    { s.contains(x) } -> std::same_as<bool>;
};

// Verifies "Parthood" or "Subset" relation (X subset of Y)
template <typename Sub, typename Super>
concept IsSubsetOf = requires {
    requires std::derived_from<Sub, Super> ||
    requires { typename Sub::is_subset_of<Super>; };
};

// CSE Pruning: Collapsing contradictory intersections to EmptySet
template <typename SetA, typename SetB>
struct Intersection {
    using type = std::conditional_t<
        HasContradictoryPredicates_v<SetA, SetB>, // Evaluated at compile-time
        EmptySet,
        InternalIntersection<SetA, SetB>
    >;
};
```

```
|};
```

Listing 25: Compile-time mereological verification using C++23 concepts.

By reifying a subset of the set-builder notation, dedekind turns the C++ translation unit into a symbolic laboratory. The programmer is free to reason in the language of Cantor, while the compiler produces the performance of a naked C++ loop.

### 3 Results

Table 10: The Dedekind Registry: Mapping C++23 Standard Primitives to Categorical Species.

| Mathematical Concept                                   | Standard C++23 Entity                 | Dedekind Concept                   | Ontological Status              |
|--------------------------------------------------------|---------------------------------------|------------------------------------|---------------------------------|
| <i>Boundary Objects (Limits)</i>                       |                                       |                                    |                                 |
| Terminal Object (1)                                    | <code>std::monostate</code>           | <code>IsTerminalObject</code>      | <b>Total</b> (Singleton)        |
| Initial Object (0)                                     | <code>std::nullptr_t</code>           | <code>IsInitialObject</code>       | <b>Total</b> (Unreachable)      |
| <i>Morphic Foundations</i>                             |                                       |                                    |                                 |
| Product ( $A \times B$ )                               | <code>std::pair&lt;A, B&gt;</code>    | <code>IsProduct</code>             | Universal Projection            |
| Coproduct ( $A + B$ )                                  | <code>std::variant&lt;A, B&gt;</code> | <code>IsCoproduct</code>           | Rule of Choice                  |
| <i>Algebraic Species (Total)</i>                       |                                       |                                    |                                 |
| Modular Group ( $\mathbb{Z}/2^n\mathbb{Z}, +$ )        | <code>unsigned int, std::plus</code>  | <code>IsAbelianGroup</code>        | <b>Total</b> (Closed)           |
| Boolean Rig ( $\mathbb{B}, \vee, \wedge$ )             | <code>bool,   , &amp;&amp;</code>     | <code>IsRig</code>                 | <b>Total</b> (Idempotent)       |
| Modular Ring ( $\mathbb{Z}/2^n\mathbb{Z}, +, \times$ ) | <code>unsigned int, +, *</code>       | <code>IsRing</code>                | <b>Total</b> (Distributive)     |
| <i>Hazardous Species (Rejected)</i>                    |                                       |                                    |                                 |
| Hazardous Magma                                        | <code>int, std::plus</code>           | <code>IsPointed</code>             | <b>Hazard</b> (UB Overflow)     |
| Floating Point Magma                                   | <code>double, std::plus</code>        | <code>IsMagma</code>               | <b>Hazard</b> (Non-associative) |
| <i>Relational Lattices</i>                             |                                       |                                    |                                 |
| Distributive Lattice                                   | <code>bool, std::logical_or</code>    | <code>IsDistributiveLattice</code> | <b>Total</b>                    |
| Pre-order ( $\sqsubseteq$ )                            | <code>operator&lt;=</code>            | <code>IsPartOf</code>              | Mereological Basis              |

#### 3.1 The Registry of Formal Correspondences

To conclude the implementation hierarchy, we provide a formal mapping between the abstract nomenclature of the mathematical literature and the concrete syntax of the C++23 language. In this structuralist posture, the “identity” of a species is not found in its nominal inheritance, but in its ability to satisfy the mereological and algebraic requirements dictated by the **Dedekind Atlas** (see Table 11).

By anchoring established axioms directly into the standard operator set—where `operator<=` serves as the formal proof of Parthood ( $\sqsubseteq$ ) and the “Highway” `operator»` represents Kleisli composition—we ensure that the resulting system is not merely a simulation of a mathematical model, but a direct execution of its laws. This registry serves as the final “Seal of Truth” for the implementation, providing the compiler with the semantic map required to navigate the ontology.

Table 11: The Dedekind Semantic Mapping: A Registry of Formal Correspondences

| Mathematical Concept        | Symbol                | C++ Operator / Type                          | C++ Concept                     | Basis                              |
|-----------------------------|-----------------------|----------------------------------------------|---------------------------------|------------------------------------|
| <i>Category Theory</i>      |                       |                                              |                                 |                                    |
| Kleisli Extension           | $(T, \eta, \gg=)$     | <code>operator&gt;&gt;=</code>               | <code>IsKleisliExtension</code> | <code>IsAssociative</code>         |
| Kleisli Extension           | $\gg=$                | <code>operator&gt;&gt;=</code>               | <code>IsKleisliExtension</code> | <code>IsAssociative</code>         |
| Kleisli Composition         | $\circ_K$             | <code>operator&gt;&gt;</code>                | <code>IsMonad</code>            | <code>IsAssociative</code>         |
| Co-Kleisli Extension        | $\gg_{\circ}$         | <code>operator&lt;&lt;=</code>               | <code>IsComonad</code>          | <code>IsAssociative</code>         |
| Morphic Injection (Unit)    | $\eta$                | <code>into&lt;F&gt; / singleton</code>       | <code>IsMonad</code>            | $\eta : T \rightarrow F(T)$        |
| Morphic Extraction (Counit) | $\varepsilon$         | <code>extract&lt;F&gt;</code>                | <code>IsComonad</code>          | $\varepsilon : F(T) \rightarrow T$ |
| Characteristic Morphism     | $\chi$                | <code>operator()</code>                      | <code>IsCharacteristic</code>   |                                    |
| Co-Kleisli Composition      | $\circ_{CoK}$         | <code>operator&lt;&lt;</code>                | <code>IsComonad</code>          |                                    |
| Counit (Extraction)         | $\varepsilon$         | <code>extract&lt;F&gt;</code>                | <code>IsComonad</code>          |                                    |
| <i>Mereology</i>            |                       |                                              |                                 |                                    |
| Proper Parthood             | $\chi$                | <code>operator()</code>                      | <code>IsProperPart</code>       | <code>IsCharacteristic</code>      |
| Parthood (Pre-order)        | $\sqsubseteq$         | <code>operator&lt;=</code>                   | <code>IsPartOf</code>           |                                    |
| Mereological Sum (Join)     | $\sqcup$              | <code>operator </code>                       | <code>IsJoinSemilattice</code>  |                                    |
| Mereological Product (Meet) | $\sqcap$              | <code>operator&amp;</code>                   | <code>IsMeetSemilattice</code>  |                                    |
| Universal Complement        | $\neg$                | <code>operator!</code>                       | <code>IsComplemented</code>     |                                    |
| Initial Object (Empty)      | $\emptyset$           | <code><math>\emptyset&lt;T&gt;</math></code> | <code>IsInitialObject</code>    |                                    |
| Terminal Object (Universal) | $\Omega$              | <code><math>\Omega&lt;T&gt;</math></code>    | <code>IsTerminalObject</code>   |                                    |
| <i>Set Theory</i>           |                       |                                              |                                 |                                    |
| Membership                  | $\in$                 | <code>operator()</code>                      | <code>IsSet</code>              | <code>IsProperPart</code>          |
| Subset                      | $\subseteq$           | <code>operator&lt;=</code>                   | <code>IsSubset</code>           | <code>IsPartOf</code>              |
| Countable Infinity          | $\aleph_0$            | <code><math>\aleph_0</math></code>           | <code>IsCardinality</code>      |                                    |
| Continuum Cardinality       | $\beth_1$             | <code><math>\beth_1</math></code>            | <code>IsCardinality</code>      |                                    |
| <i>Order Theory</i>         |                       |                                              |                                 |                                    |
| Pre-order                   | $\sqsubseteq$         | <code>operator&lt;=</code>                   | <code>IsPartiallyOrdered</code> | <code>IsPartOf</code>              |
| <i>Algebra</i>              |                       |                                              |                                 |                                    |
| Structural Origin (Zero)    | 0                     | <code>T::origin()</code>                     | <code>IsPointed</code>          |                                    |
| Additive Composition        | +                     | <code>operator+</code>                       | <code>IsGroup</code>            |                                    |
| Additive Inverse            | -                     | <code>operator-</code>                       | <code>IsGroup</code>            |                                    |
| Multiplicative Composition  | $\times$              | <code>operator*</code>                       | <code>IsRing</code>             |                                    |
| Multiplicative Inverse      | $/, ^{-1}$            | <code>operator/</code>                       | <code>IsField</code>            |                                    |
| Associativity               | $(a \circ b) \circ c$ | <code>is_associative_v</code>                | <code>IsSemigroupoid</code>     |                                    |
| Idempotency                 | $a \circ a = a$       | <code>is_idempotent_v</code>                 | <code>IsIdempotent</code>       |                                    |
| Magnitude / Count           | $ S $                 | <code>s.size() / s.cardinality()</code>      | <code>IsCardinality</code>      |                                    |
| Supremum / Infimum          | $\sup, \inf$          | <code>s.supremum(), s.infimum()</code>       | <code>HasExtrema</code>         |                                    |

### 3.2 order Module

### 3.3 topology Module

### 3.4 sequences Module

### 3.5 algebra Module

Needed to push back repeatedly on attempts to go the convenient ‘double’ route instead of spelling things out as semi-modules, semi-rings etc. i.e. reduced-strength algebraic structures. Of course, ultimately everything is a complex number from this perspective just like everything is a string for front-end developers.

### 3.6 morphologies Module

### 3.7 geometry Module

### 3.8 numbers Module

### 3.9 analysis Module

The physical build order of the `dedekind.ontology` module mirrors this conceptual hierarchy. By enforcing a linear dependency graph—where `:mereology` is compiled before `:algebra`—we ensure that the compiler’s proof-search for a Ring’s identity element is grounded in the prior verification of its mereological subobjects. This hierarchy transforms the C++ module mapper into a directed acyclic graph of mathematical truth.

## 4 Discussion

However, reifying a formal ontology in a statically-typed **systems** language presents significant challenges compared to "pure" functional languages like Haskell [39] or Agda [40], which have native support for higher-kinded types, dependent types and first-class functions. Challenges encountered include:

- **Template Template Parameters:** C++ is notoriously bad at true Category Theory because it lacks higher-kinded Types (HKTs). As a result `Functor<F>` or `Monad<F>` are defined in terms of `template template`, which have significant edge cases (e.g., they don’t play well with aliases or multiple template arguments).
- **The Lambda Hurdle (Opaqueness):** Unlike the first-class, inspectable functions in the ML family, C++ lambdas are unique, anonymous types. They do not inherently expose their domain or codomain to the type system. A morphism’s domain and codomain must be inferred at the call site or explicitly provided by the programmer. Although if used as *Non-Type Template Parameters* (NTTP), they can be "inspected" them via `concept` checks or by wrapping them in a type that carries its own signature metadata.
- **Decidability vs. Instantiation:** While dependent-type systems use proof-search, C++ relies on template instantiation. This introduces the risk of "Late Detection," where a mathematical violation is only discovered when a set is materialized, rather than when it is defined. This can be mitigated by forcing evaluations into `consteval` and `constexpr` contexts.
- **The Performance-Rigidity Tradeoff:** High-level abstractions in C++ often incur a "template tax" in compile times. We utilize C++23 **Module Partitions** to create a directed acyclic graph of truth, caching structural invariants to ensure that formal rigor does not collapse the build pipeline.

- **The Specialization Constraint:** Unlike the unified pattern matching of functional languages, C++ relies on Partial Template Specialization to resolve structural properties. This creates a "Rigidity Challenge": the compiler cannot readily deduce that a property of  $A$  also applies to  $A \cup B$  without explicit boilerplate. We address this by using `concept ... requires` as logical compile-time predicate, thereby simulating the polymorphic reach of a true dependent-type system.
- **The Precedence-Associativity Constraint:** A significant friction point in embedding categorical notation within the C++ grammar is the language's fixed set of operator candidates and their immutable precedence rules. While `operator>>=` is the "nominal" choice for Monadic Bind ( $\gg=$ ), the ISO standard defines all compound assignment operators with right-to-left associativity. Consequently, a standard monadic chain such as `m >>= f >>= g` is parsed as `m >>= (f >>= g)`, which violates the directional vector of the Kleisli triple.

Rather than yielding to this syntactic prejudice, `dedekind` introduces the **Highway Notation** (`operator>>`). This is not a mere workaround for the "Bind" associativity; it is a deliberate architectural choice to prioritize the left-to-right "push" of the structural pipeline. By utilizing the left-associativity of the shift operator, the Highway reifies the categorical flow as a first-class visual entity. This ensures that the code's "scent" remains isomorphic to the mathematical intent, even when the language's formal grammar resists the composition.

## 5 Conclusion

This paper has demonstrated the feasibility of embedding formal mathematical structuralism within the C++23 type system through the `dedekind` library. By transitioning from the traditional object-oriented paradigm of nominal inheritance toward a model of *structural transparency*, we have illustrated how the modern compiler can be utilized as a deterministic constraint satisfaction engine. This approach allows for the high-level representation of complex mathematical species—such as intensional set-builder notation and the Dedekind-complete continuum—without the runtime performance penalties historically associated with abstract symbolic modeling.

Our investigation suggests that the Clang/LLVM toolchain, while not a dedicated theorem prover, is capable of implementing a "rethought" set theory where categorical axioms provide the semantic context for aggressive machine-level optimization. By hoisting mereological relations into the type signature, we enable the backend to perform *structural pruning*: collapsing logical contradictions and optimizing intersections into zero-instruction sequences that maintain the efficiency of hand-tuned assembly.

The stratification of these formalisms into verified module partitions provides a scalable pathway for managing the "template tax," ensuring that the rigors of formal logic do not destabilize the build pipeline. By embedding the construction of the Continuum directly into the translation unit, we have moved toward a nascent paradigm of **Axiomatic Systems Programming**. In this model, the build process acts as a physical guardrail for mathematical truth, ensuring that high-level algebraic invariants and low-level hardware sympathy are no longer dualities, but a unified, verified reality.

Future work will extend this mereological framework to the upcoming linear algebra capabilities of the C++ standard library [41], expressly facilitating exact computations across semimodules and *reduced* (booleans, `char`, `unsigned int`) or *extended* (complex, dual, quaternions) number systems. By anchoring these structures in the `dedekind` atlas, we aim to provide a unified algebraic substrate that transcends the binary-limitations of traditional hardware types, enabling a truly symbolic systems programming environment.

## References

- [1] Adam Paszke, Sam Gross, Francisco Massa, et al. “PyTorch: An Imperative Style, High-Performance Deep Learning Library”. In: *Advances in Neural Information Processing Systems 32*. Ed. by H. Wallach et al. 2019, pp. 8024–8035. URL: <https://neurips.cc>.
- [2] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, et al. “Array programming with NumPy”. In: *Nature* 585.7825 (2020), pp. 357–362. DOI: 10.1038/s41586-020-2649-2.
- [3] Wenzel Jakob. *nanobind: tiny and efficient C++/Python bindings*. <https://github.com/wjakob/nanobind>. 2022.
- [4] ISO/IEC JTC 1/SC 22. *ISO/IEC 14882:2024: Programming languages — C++*. Commonly referred to as C++23. International Organization for Standardization. Geneva, Switzerland, 2024. URL: <https://www.iso.org>.
- [5] Ivan Čukić. *Functional Programming in C++*. Manning Publications, 2018. ISBN: 9781617293818.
- [6] Eugenio Moggi. “Notions of computation and monads”. In: *Information and computation* 93.1 (1991), pp. 55–92.
- [7] Steve Awodey. “An Answer to Hellman’s Question: Does Category Theory Provide a Framework for Mathematical Structuralism?” In: *Philosophia Mathematica* 12.1 (2004), pp. 54–64.
- [8] André Joyal. “Une théorie combinatoire des séries formelles”. In: *Advances in Mathematics* 42.1 (1981), pp. 1–82. DOI: 10.1016/0001-8708(81)90052-9.
- [9] André Joyal. “Fonctions analytiques et espèces de structures”. In: *Combinatoire énumérative*. Springer, 1986, pp. 126–159. DOI: 10.1007/BFb0072514.
- [10] Colin McLarty. “Numbers can be just what they have to”. In: *Noûs* 27.4 (1993), pp. 487–498.
- [11] John Baez, Urs Schreiber, and David Corfield. *The n-Category Café: A group blog on math, physics and philosophy*. <https://utexas.edu>. 2006–present.
- [12] IEEE. *IEEE Standard for Floating-Point Arithmetic*. Provides the formal specification for floating-point formats and operations. 2019. DOI: 10.1109/IEEESTD.2019.8766229.
- [13] Martin Thompson. *Mechanical Sympathy*. Mechanical Sympathy Blog. 2011. URL: <https://mechanical-sympathy.blogspot.com/2011/07/why-mechanical-sympathy.html> (visited on 03/31/2026).
- [14] Koen Claessen and John Hughes. “QuickCheck: a lightweight tool for random testing of Haskell programs”. In: *Proceedings of the 5th ACM SIGPLAN International Conference on Functional Programming (ICFP)*. ACM, 2000, pp. 268–279.
- [15] Bryan O’Sullivan, John Goerzen, and Don Stewart. *Real World Haskell*. O’Reilly Media, 2008.
- [16] Philip Wadler. “Theorems for free!” In: *Conference on Functional Programming Languages and Computer Architecture*. ACM. 1989, pp. 347–359.
- [17] William Shakespeare. *The Most Excellent and Lamentable Tragedy of Romeo and Juliet*. Act 2, Scene 2. John Danter, 1597.
- [18] Bartosz Milewski. *Category Theory for Programmers*. Compiled from the original blog series at <https://bartoszmilewski.com>. Blurb, Incorporated, 2019. ISBN: 9780464243878.
- [19] Colin McLarty. *Elementary Categories, Elementary Toposes*. Clarendon Press, 1992.

- [20] Helena Rasiowa. *An Algebraic Approach to Non-Classical Logics*. Foundational for the Subobject Classifier ( $\Omega$ ) and the algebraic treatment of truth-values as objects in a Topos. North-Holland Publishing Company, 1974.
- [21] F. William Lawvere. “An elementary theory of the category of sets”. In: *Proceedings of the National Academy of Sciences* 52.6 (1964), pp. 1506–1511.
- [22] Kazimierz Ajdukiewicz. *Język i poznanie: Wybór pism z lat 1920-1939*. Vol. 1. Foundational text for the Polish School of Logic and Semantic Epistemology. Państwowe Wydawnictwo Naukowe, 1985.
- [23] J. R. B. Cockett and Stephen Lack. “Restriction categories I: categories of partial maps”. In: *Theoretical Computer Science* 270.1-2 (2002). Formalizes the distinction between total and partial morphisms, providing the foundation for the `:total partition`., pp. 223–259.
- [24] Richard Hubert Bruck. *A Survey of Binary Systems*. Ergebnisse der Mathematik und ihrer Grenzgebiete. The definitive reference for the structural hierarchy of Quasigroups and Loops, justifying the decoupling of associativity from invertibility. Springer-Verlag, 1971.
- [25] Carl Pomerance. “The role of smooth numbers in number theoretic algorithms”. In: *Proceedings of the International Congress of Mathematicians*. Provides the number-theoretic context for ‘Lipschitz boundaries’ and the rejection of algebraic closure in finite machine integers. 2001.
- [26] Urs Stammach. *Lineare Algebra*. A foundational text emphasizing the structural and algebraic aspects of linear algebra. Teubner, 1994.
- [27] Paul R. Halmos. *Finite-Dimensional Vector Spaces*. The University Series in Undergraduate Mathematics. A seminal text advocating for a coordinate-free, structuralist approach to linear transformations, aligning with the concept-based dispatch in `dedekind`. D. Van Nostrand Company, 1958.
- [28] Heinrich Kleisli. “Every standard construction is induced by a pair of adjoint functors”. In: *Proceedings of the American Mathematical Society* 16.3 (1965), pp. 544–546. DOI: 10.2307/2034658.
- [29] Bartosz Milewski. *Category Theory for Programmers*. An excellent bridge between C++ type systems and categorical structures. Blending Science, 2018.
- [30] nLab authors. *Kleisli category*. <https://ncatlab.org>. Accessed: 2026-04-04. 2026.
- [31] F. William Lawvere and Robert Rosebrugh. *Sets for Mathematics*. Discusses the construction of the continuum and Dedekind cuts within a categorical framework. Cambridge University Press, 2003.
- [32] F William Lawvere. “Metric spaces, generalized logic, and closed categories”. In: *Rendiconti del seminario matematico e fisico di Milano* 43.1 (1973), pp. 135–166.
- [33] Francis Borceux and Françoise Dejean. “Cauchy completion in category theory”. In: *Cahiers de topologie et géométrie différentielle catégoriques* 27.2 (1986), pp. 133–146.
- [34] Thomas Mormann. “Structural Mereology and Statistics”. In: *Logic and Logical Philosophy* 22.4 (2013), pp. 427–453.
- [35] F William Lawvere. “An elementary theory of the category of sets (long version) with commentary”. In: *Reprints in Theory and Applications of Categories* 11 (2005), pp. 1–35.
- [36] Joachim Lambek and Philip J Scott. *Introduction to Higher-Order Categorical Logic*. Cambridge University Press, 1988.
- [37] nLab authors. *structural set theory*. [ncatlab.org/nlab/show/structural+set+theory](https://ncatlab.org/nlab/show/structural+set+theory). 2024. URL: <https://ncatlab.org>.



- [38] Jean-Pierre Marquis. “The Structuralist Mathematical Style: Bourbaki as a Case Study”. In: *The Architecture of Modern Mathematics*. Oxford University Press, 2022.
- [39] The GHC Development Team. *The Glasgow Haskell Compiler User’s Guide, Edition 2024*. GHC 2024 Language Edition. 2024. URL: [https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/control.html](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/control.html).
- [40] Ulf Norell. “Towards a practical programming language based on dependent type theory”. PhD thesis. Chalmers University of Technology, 2007. URL: <http://www.cse.chalmers.se>.
- [41] Mark Hoemmen et al. *P1673R13: A free function linear algebra interface based on the BLAS*. Published Proposal. Approved for inclusion in the C++26 Standard Library. ISO/IEC JTC1/SC22/WG21, Nov. 2023. URL: <https://wg21.link>.